

Peer to Peer Networks

P2P Application Examples

- File Sharing Applications (esp.Mp3 music)
 - Napster, Gnutella, KaZaA, eDonkey, eMule, KAD, ...
- Video downloads, video-on-demand (VoD), large-scale file distribution (e.g., software updates)
 - BitTorrent & variants, ...
- Skype (P2P VoIP & video conferencing)
- (real-time/near-real time) video broadcasting, video streaming, ...
 - pplive, qqlive,
- Technology Applications
 - Dynamo: Amazon storage substrate
 - Akamai Netsession: P2P streaming client
 - WebRTC for browser-browser video streaming

Peer to Peer Networks

- Reducing load of servers
- Faster distribution of content
- Scalable content distribution
- Low barrier to deployment
- Organic growth (self-scalable)
- Resilience to attacks and failures
- Resources grow with size of network
- Churn of peers makes network unstable

Locating the Relevant Peers


- Three main approaches
 - Central directory (Napster)
 - Query flooding (Gnutella)
 - Hierarchical overlay (Kazaa, modern Gnutella)
- Design goals
 - Scalability
 - Simplicity
 - Robustness
 - Plausible deniability

P2P (Application) Model

- Each user stores a subset of files (content)
- Each user has access (can download) files from all users in the system

Key Challenges in “pure” peer-to-peer model

- How to locate your peer & find what you want?
- Need some kind of “directory” or “look-up” service



- centralized
- distributed, using a hierarchical structure
- distributed, using a flat structure
- distributed, with no structure (“flooding” based)
- distributed, using a “hybrid” structured/unstructured approach

5

Other Challenges

Technical:

- Scale: up to hundred of thousands or millions of machines
- Churn: machines can come and go any time

Social, economic & legal:

- Incentive Issues: free-rider problem
 - Vast majority of users are free-riders
 - most share no files and answer no queries
 - A few individuals contributing to the “public good”
- Copyrighted content and piracy
- Trust & security issues

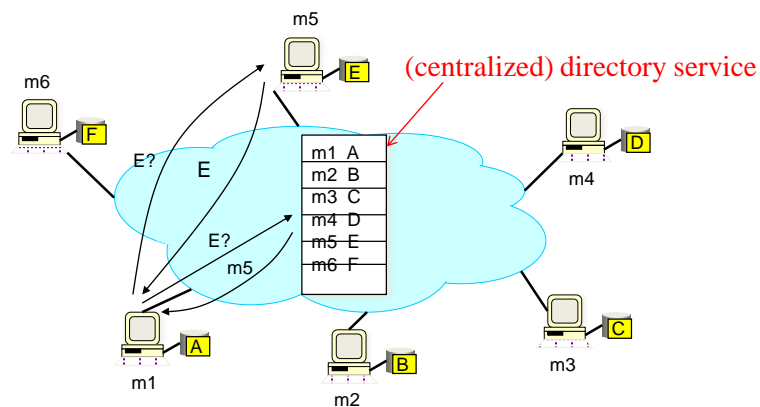
6

Unstructured P2P Applications

- Napster
 - a *centralized* directory service
 - peers directly download from other peers
- Gnutella
 - fully distributed directory service
 - discover & maintain neighbors, *ad hoc* topology
 - flood & forward queries to neighbors (with bounded hops)
- Skype
 - exploit heterogeneity, certain peers as “super nodes”
 - two-tier hierarchy: when join, contact a super-node
 - smart query flooding
 - peer may fetch data from multiple peers at once
- Pros and Cons of each approach?

7


Napster Architecture: An Illustration



8

Gnutella


- Ad-hoc topology
- Queries are flooded for bounded number of hops
- No guarantees on recall



9

KaZaA: Exploiting Heterogeneity


- Each peer is either a group leader or assigned to a group leader (supernode)
 - TCP connection between peer and its group leader
 - TCP connections between some pairs of group leaders
- Group leader tracks the content in all its children
- Q: how to select supernodes?



10

Voice-over-IP: Skype


- proprietary application-layer protocol (inferred via reverse engineering)
 - encrypted msgs
- P2P components:
 - **clients:** Skype peers connect directly to each other for VoIP call
 - **super nodes (SN):** Skype peers with special functions
 - **overlay network:** among SNs to locate SCs
 - **login server**



P2P voice-over-IP: Skype


Skype client operation:

1. joins Skype network by contacting SN (IP address cached) using TCP
2. logs-in (username, password) to centralized Skype login server
3. obtains IP address for callee from SN, SN overlay
 - or client buddy list
4. initiate call directly to callee



Skype: peers as relays

- **problem:** both Alice, Bob are behind “NATs”
 - NAT prevents outside peer from initiating connection to insider peer
 - inside peer *can* initiate connection to outside
- **relay solution:** Alice, Bob maintain open connection to their SNs
 - Alice signals her SN to connect to Bob
 - Alice’s SN connects to Bob’s SN
 - Bob’s SN connects to Bob over open connection Bob initially initiated to his SN




P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks


tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file

Alice arrives ...
... obtains list of peers from tracker
... and begins exchanging file chunks with peers in torrent




BitTorrent: Overall Architecture



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

Requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

Sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4


BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Spotify

- Uses BT as basic protocol
 - Uses server for first 15s
 - Tries to find peers and download from them
 - Only 8.8% of bytes come from servers
- When 30s left
 - Starts searching for next tr
 - Uses sever with 10s to go i no peers found



19

BitTorrent & Video Distribution

- Designed for large file (e.g., video) downloads
 - esp. for popular content, e.g. flash crowds
- Focused on efficient *fetching*, not search
 - Distribute same file to many peers
 - Single publisher, many downloaders
- Divide large file into many pieces
 - Replicate different pieces on different peers
 - A peer with a complete piece can trade with other peers
- Allows simultaneous downloading
 - Retrieving different parts of the file from different peers at the same time
 - Fetch rarest piece first
- Also includes mechanisms for preventing “free riding”


20

Structured P2P Networks

- Introduce a **structured** logical topology
- Abstraction: a **distributed hash table** data structure
 - `put(key, object); get (key)`
 - `key`: identifier of an object
 - `object` can be anything: a data item, a node (host), a document/file, pointer to a file, ...
- Design Goals: guarantee on recall
 - i.e., ensure that an item (file) identified is always found
 - Also scale to hundreds of thousands of (or more) nodes
 - handle rapid join/leave and failure of nodes
- Proposals
 - Chord, CAN, Kademlia, Pastry, Tapestry, etc

21

Distributed hash table



DHT provides the information look up service for P2P applications.

- Nodes uniformly distributed across *key* space
- Nodes form an *overlay* network
- Nodes maintain list of neighbors in routing table
- Decoupled from physical network topology

22

Key Ideas (Concepts & Features)

- Keys and node ids map to the same “flat” id space
 - node ids are thus also (special) keys!
- Management (organization, storage, lookup, etc) of keys using consistent hashing
 - distributed, maintained by all nodes in the network
- (Logical) distance defined on the id space: structured!
 - different DHTs use different distances/structures
- Look-up/Routing Tables (“finger table” in Chord)
 - each node typically maintains $O(\log n)$ routing entries
 - organizing using structured id space: more information about nodes closer-by; less about nodes farther away
- Bootstrap, handling node joins/leaves or failures
 - when node joins: needs to know at least one node in the system
- Robustness/resiliency through redundancy

23

DHT identifiers

- assign integer identifier to each peer in range $[0, 2^n - 1]$ for some n .
 - each identifier represented by n bits.
- require each key to be an integer in same range
- to get integer key, hash original key
 - e.g., key = hash(“Led Zeppelin IV”)
 - this is why its is referred to as a distributed “hash ” table


Consistent Hashing

- Every key is hashed to generate a number in a circular range
- Every node/replica assigned an ID in the same space
- A key is stored at the first N nodes which succeed the hash of the key in the circular ring
 - Called “preference list” of the key
- First node on the list is the coordinator for the key
 - Get/put operations at all nodes managed by coordinator
- For better load balancing, every node is treated as multiple virtual nodes, assigned many positions on list
 - Preference list will contain N distinct physical nodes

Consistent Hashing Properties

- **Load balance:** all nodes receive roughly the same number of keys
- For N nodes and K keys, with high probability
 - each node holds at most $(1+\epsilon)K/N$ keys
 - (provided that K is large enough compared to N)


Circular DHT (I)




- each peer *only* aware of immediate successor and predecessor.
- “overlay network”

Circular DHT (I)

$O(N)$ messages
on average to resolve
query, when there
are N peers



Circular DHT with finger table



- each peer keeps track of IP addresses of predecessor, successor, short cuts.
- reduced from 6 to 2 messages.
- possible to design shortcuts so $O(\log N)$ neighbors, $O(\log N)$ messages in query