

**CS343: Operating System**

# **Synchronization**

**Lect19 : 12th Sept 2023**

**Dr. A. Sahu**

**Dept of Comp. Sc. & Engg.**

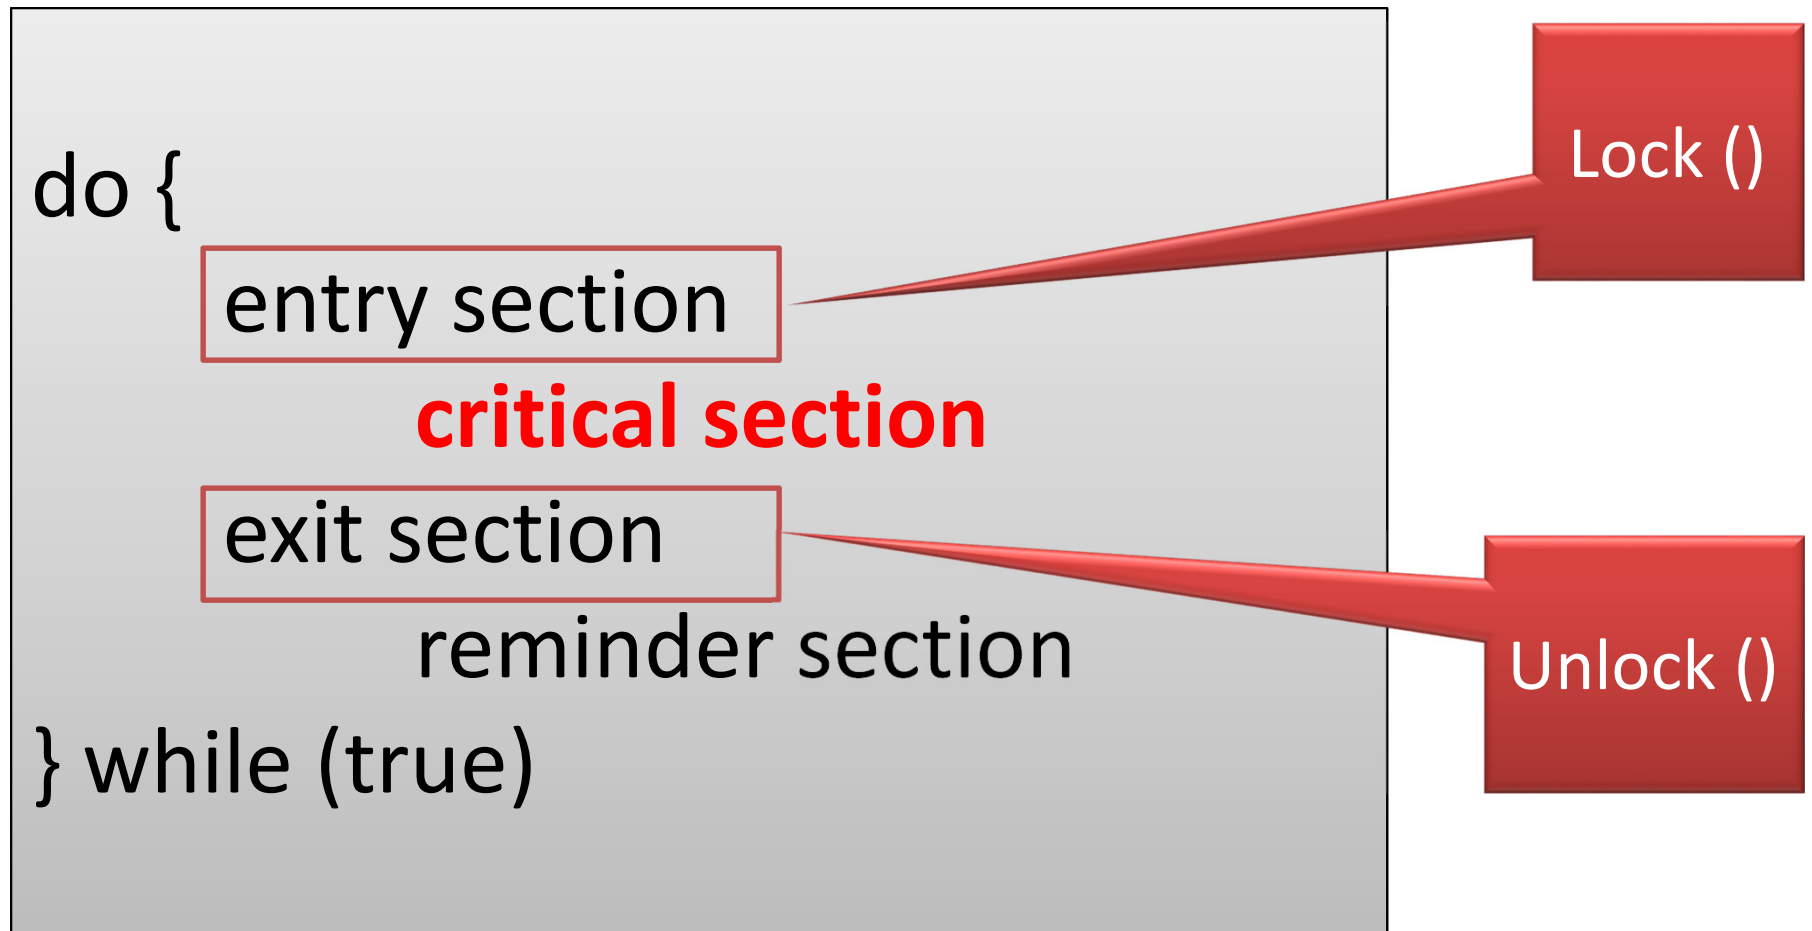
**Indian Institute of Technology Guwahati**

# Outline

- Synchronization
  - Critical Section Problem
- Solution to CS Problems
  - Two Threads Solutions: Peterson's Solution
  - N Thread Solutions: Filter and Bakery Algorithms
- Sync Hardware
  - CAS, TAS, LL-LC, BackupLock

# Critical Section

- General structure of process  $P_i$



# Peterson's Algorithm: Combine LockOne and LockTwo

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```

Announce I'm interested

Defer to other

No longer interested

Wait while other interested & I'm the victim

## Peterson's Lock: Lock 3

- Satisfy Mutual Exclusion
- Satisfy Deadlock Free
- Satisfy Starvation Free


—Proof .....

# Nthread Synchronization

# Filter Algorithm for $n$ Threads

There are  $n-1$  “waiting rooms” called levels

- At each level
  - At least one enters level
  - At least one blocked if  
many try
- Only one thread makes it through




# Filter

```
class FilterLock {  
    int level[n]; // level[i] for thread i  
    int victim[n]; // victim[L] for level L  
    public FilterInit(int n) {  
        for(int i=1; i<n; i++)  
            level[i]=0;  
        ...  
    }  
}
```



# Filter




Thread 2 at level 4

# Filter

```
class FilterLock {  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k!=i level[k]>=L) &&  
                    victim[L]==i) {};  
        }  
    }  
    public void unlock() { level[i]=0; }  
}
```

# Claim: Mutex

- Start at level  $L=0$
- At most  $n-L$  threads enter level  $L$
- Mutual exclusion at level  $L=n-1$



# No Starvation

- Filter Lock satisfies properties:
  - Just like Peterson Alg at any level
  - So no one starves
- But what about fairness?
  - Threads can be overtaken by others**

# Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- If A starts before B, then A enters before B?
- But what does “start” mean?
- Need to adjust definitions ....

# Bakery Algorithm

- Similar to Bakery Shop
- Provides First-Come-First-Served
- How?
  - Take a “number”
  - Wait until lower numbers have been served
- Lexicographic order
  - $(a,i) > (b,j)$ 
    - If  $a > b$ , or  $a = b$  and  $i > j$

# Bakery Algorithm


```
class BakeryLock {  
    bool flag[n];  
    int label[n];  
    public BakeryLockInit(int n) {  
        for(int i = 0; i < n; i++) {  
            flag[i]= false; label[i] = 0;  
        }  
    }  
    ...  
}
```

# Bakery Algorithm

```
class Bakery Lock {
```

```
    bool flag[n];
```

```
    int label[n];
```



CS



# Bakery Algorithm

```
class BakeryLock {  
    public void lock() {  
        flag[i]=true;  
        label[i]=max(label[0], ..., label[n-1])+1;  
        while (  $\exists k$  flag[k] &&  
                label[i] > label[k]) {}  
    }  
}
```

# Mutual Exclusion

- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
  - $\text{flag}[A]$  is *false*, or
  - $\text{label}[A] > \text{label}[B]$

# No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)

```
label[i] = max(label[0], ..., label[n-1]) + 1;
```

# Synchronization Hardware

# Synchronization Hardware

- Many systems provide hardware support for Sync.
- All solutions below based on idea of **locking**
  - Protecting critical regions via locks
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - Operating systems using this not broadly scalable

# Synchronization Hardware

- Multiprocessor–disable interrupts
  - Generally too inefficient on multiprocessor systems
    - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - **Atomic** = non-interruptible
    - Either test memory word and set value
    - Or swap contents of two memory words

# Atomic Sync Instructions

- Test and Set (**TAS**)
- Compare and Swap (**CAS**)
- Exchange (**XCHG**)
- Fetch and Increment (**FAI**)
- **How to provide these in MP/MC?**
  - Load Locked & Store Conditional (**LL-SC**)

# test\_and\_set Instruction : TAS

//Definition:

```
boolean test_and_set (boolean *target) {  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to "TRUE".



# Solution using test\_and\_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

# CAS Instruction

```
int CAS(int *value,  
        int expected, int new_value) {  
  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;  
    return temp;  
}
```

- Executed atomically
- Returns the original value of passed parameter “value”
- Set the variable “value” the value of the passed parameter “new\_value” but only if “value” == “expected”. That is, the swap takes place only under this condition.

# Sync Solution using CAS

- Shared integer “lock” initialized to 0;

```
do {  
    while (CAS(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
  
} while (true);
```

# Synchronization Instruction

- Hardware primitive for atomic read+write is required e.g.
  - Exchange, (XCHG)
  - Test & Set (TAS)
    - // test for unlock (0) then set the lock (1)
  - Fetch & Increment (FAI)

# Spin Lock with Exchange Instr.

Lock: 0 indicates free and 1 indicates locked

Code to lock X :

r2 = 1

**SPINing**

Try to test and acquire the lock in a tight loop

```
lockit:    r2  $\leftrightarrow$  X //atomic exchange  
          if(r2 $\neq$ 0) goto lockit //already locked
```

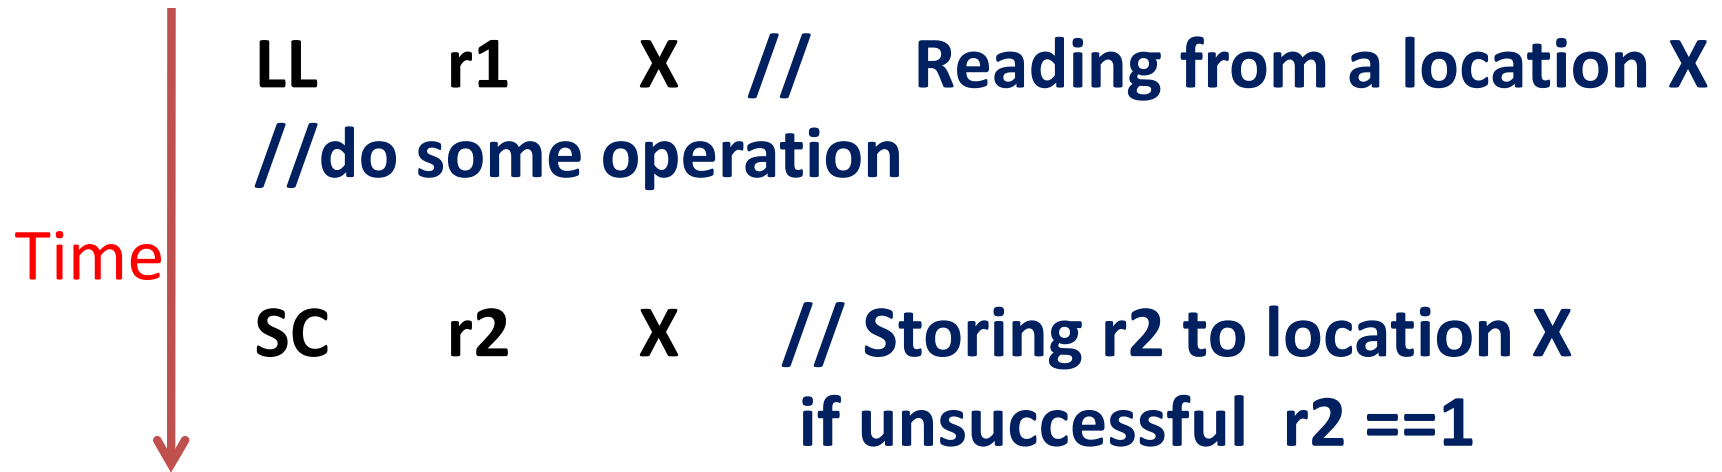
locks are cached for efficiency, coherence is used in shared multiprocessor

# Spin Lock with Exchange Instr.

## Better code to lock X :

```
lockit:    r2 = X;      // read lock
           if(r2≠0) goto lockit ; //not available
           r2 = 1;
           r2 ↔ X ;    //atomic exchange
           if(r2≠0) goto lockit ; //already locked
```

# LD Locked & ST conditional



Store will be unsuccessful if values of X is altered/changed by others processor between time of LL and time of SC you have done

Load linked/Store Conditional

# Spin Lock with LL & SC

lockit:

```
LL r2, X           //load locked
if(r2≠0) goto lockit ; //not available
r2 = 1 ;
SC r2, X           //store conditional
if(r2==1) goto lockit ; // store fails redo
```

**Spin lock with exponential back-off reduces contention**



# Atomic XCHG with LL & SC

Simpler to implement

- **Atomic exchange using LL and SC**

```
try:  r3 = r2;           //move exchange value
      LL r1, X           // load locked
      SC r3, X           //store conditional
      if(r3==1) goto try ; //store fails redo...
      r2 =r1;            //put loaded value in r2
```

# Atomic FAI with LL & SC

Simpler to implement

- Fetch & increment using LL and SC

```
try:  LL r1, X           //load locked
      r3 = r1 + 1;        //increment
      SC r3, X            //store conditional
      if(r3==1) goto try ; //store fails redo...
```