

CS343: Operating System

**Synchronization: Semaphore,
Monitors**

Lect22 : 26th Sept 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Synchronization
 - Critical Section Problem
- Sync Hardware
 - CAS, TAS, LL-LC, BackupLock
- Semaphore and Monitor
- Classical Sync Problems

Semaphore

- Semaphore: Synchronization tool
 - Provides more sophisticated ways (than Mutex)
 - For process to synchronize their activities.
- Semaphore: Abstract data type
- Semaphore **S** – integer variable
- Can only be accessed via two indivisible (atomic) operations
 - Used for controlling access, by multiple processes
 - Can be access by two atomic **Wait()** and **Signal()**

Semaphore Usage

- **Binary semaphore** – integer value can range only between 0 and 1
 - Same as a **mutex lock**
- **Counting semaphore** – integer value can range over an unrestricted domain
- Can solve various synchronization problems

Counting/Bin Semaphore

```
S=50; // Initialized to 1 for Binary  
// S =0 Locked; S>=1 Available  
void synchronized wait(S) {  
    while (S <= 0) ; // busy wait  
    S--;  
}  
  
void synchronized signal(S) {  
    S++;  
}
```

Counting Semaphore: Real life Example

- Counting Semaphores : Representation of a limited number of resources
- If a restaurant has a capacity of **50 people**
 - And nobody is there, the semaphore would be initialized to **50**



Counting Semaphore: Real life Example

- **As each person arrives at the restaurant**
 - They cause the seating capacity to decrease
 - So the semaphore in turn is decremented.
- **When the maximum capacity is reached**
 - The semaphore will be at zero
 - Nobody else will be able to enter the restaurant.
 - Instead the hopeful restaurant goers must wait until someone is done eating.
- **When a patron leaves**
 - The semaphore is incremented
 - And the resource becomes available again.

Semaphore Usage

- Consider P_1 and P_2 that require S_1 to happen before S_2

Create a semaphore “synch” initialized to 0

P1:

S_1 ;
signal(synch);

P2:

wait(synch);
 S_2 ;

- Can implement a counting semaphore S as a binary semaphore

```
wait(){while(S<=0);S--;}  
signal(){S++;}
```


Semaphore Implementation

- Guarantee that no two processes can execute
 - **wait()** and **signal()** on the same semaphore at the same time
 - **void synchronized wait(); void synchronized signal();**
- Implementation becomes the critical section problem
 - where the **wait** and **signal** code are placed in the critical section
 - Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied

Semaphore with no Busy waiting

- With each semaphore there is an associated **waiting queue**
- Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record in the list
- Two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

Semaphore with no Busy waiting


```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1



Deadlock and Starvation

- **Starvation – indefinite blocking**
 - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Solved via **priority-inheritance protocol**

Classical Problems of Synchronization

- Classical problems **used to test** newly-proposed synchronization schemes
 1. Bounded-Buffer Problem
 2. Readers and Writers Problem
 3. Dining-Philosophers Problem

Bounded-Buffer Problem

- n buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value n

Bounded Buffer: Producer

```
do {  
    ..  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ..  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```

```
wait(S) { while (S <= 0) ; Add(P,S); S--; }  
signal(S) { S++; wakeup rem(S.front); }  
// Initialization Full=0; Empty=n
```


Bounded Buffer : Consumer

```
do {  
    wait(full);  
    wait(mutex);  
    /* remove an item from buffer to next_consumed */  
    signal(mutex);  
    signal(empty);  
    /* consume the item in next consumed */  
} while (true);
```

```
wait(S) {while (S <= 0) ; Add(P,S); S--; }  
signal(S) {S++; wakeup rem(S.front);}  
// Initialization Full=0; Empty=n
```

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
 - **Readers** – only read the data set; they do *not* perform any updates
 - **Writers** – can both read and write
- Variation one: other variations.....
 - Allow multiple readers to read at the same time
 - Only one single writer can access the shared data at the same time
 - **No reader kept waiting unless writer has permission to use shared object**

Readers-Writers Problem

- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore **rw_mutex** initialized to 1
 - Semaphore **mutex** initialized to 1
 - Integer **read_count** initialized to 0

Writer Structure : RW problem

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    /* Writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

```
wait(S) {while (S <= 0) ; S--;}  
signal(S) {S++;}
```

Reader Structure : RW problem

- Multiple reader are allowed: one is reading then wait for write to complete

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    /* reading is performed */  
    wait(mutex);  
    read count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```

Readers-Writers Problem Variations


- **First variation** – no reader kept waiting unless writer has permission to use shared object
- **Second variation** – once writer is ready, it performs the write ASAP
 - Writer have very high priority
- Both may have starvation leading to even more variations

Readers-Writers Problem Variations

- ***Third variation:*** No thread shall be allowed to starve
 - Operation of obtaining a lock on the shared data will always terminate in a bounded amount of time.
 - Problem is solved on some systems by kernel providing reader-writer locks

Dining-Philosophers Problem

- Philosophers spend their lives alternating thinking and eating



Dining-Philosophers Problem

- Philosophers spend their **lives alternating thinking and eating**
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - **Need both to eat, then release both when done**
- In the case of 5 philosophers
 - Shared data : Bowl of rice,
 - Semaphore **chopstick [5]** initialized to 1

Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
    // eat  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
    // think  
} while (TRUE);
```

- What is the problem with this algorithm?

Problem in Dining-Philosophers Algorithm

- May be deadlock
 - Every one is Holding one Fork and requesting for other
 - Form a circular wait

Problem in Dining-Philosophers Algorithm

- Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table.
 - 5 chop sticks, 4 people: pigeon-hole principle at least one can easily acquire 2 chop sticks , so no deadlock
 - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
 - Allow both or none, so only two person can get chance and there will not be any deadlock

Problem in Dining-Philosophers Algorithm

- Deadlock handling: Use an asymmetric solution
 - An odd-numbered philosopher picks up first the left chopstick and then the right chopstick. **Left then Right**
 - Even-numbered philosopher picks up first the right chopstick and then the left chopstick. **Right then Left**
 - As neighbor are different and circular fashion : They will allow you to pickup (if all try at same time)