

CS343: Operating System

Memory Management

Lect30 : 16th Oct 2023

Dr. A. Sahu

Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

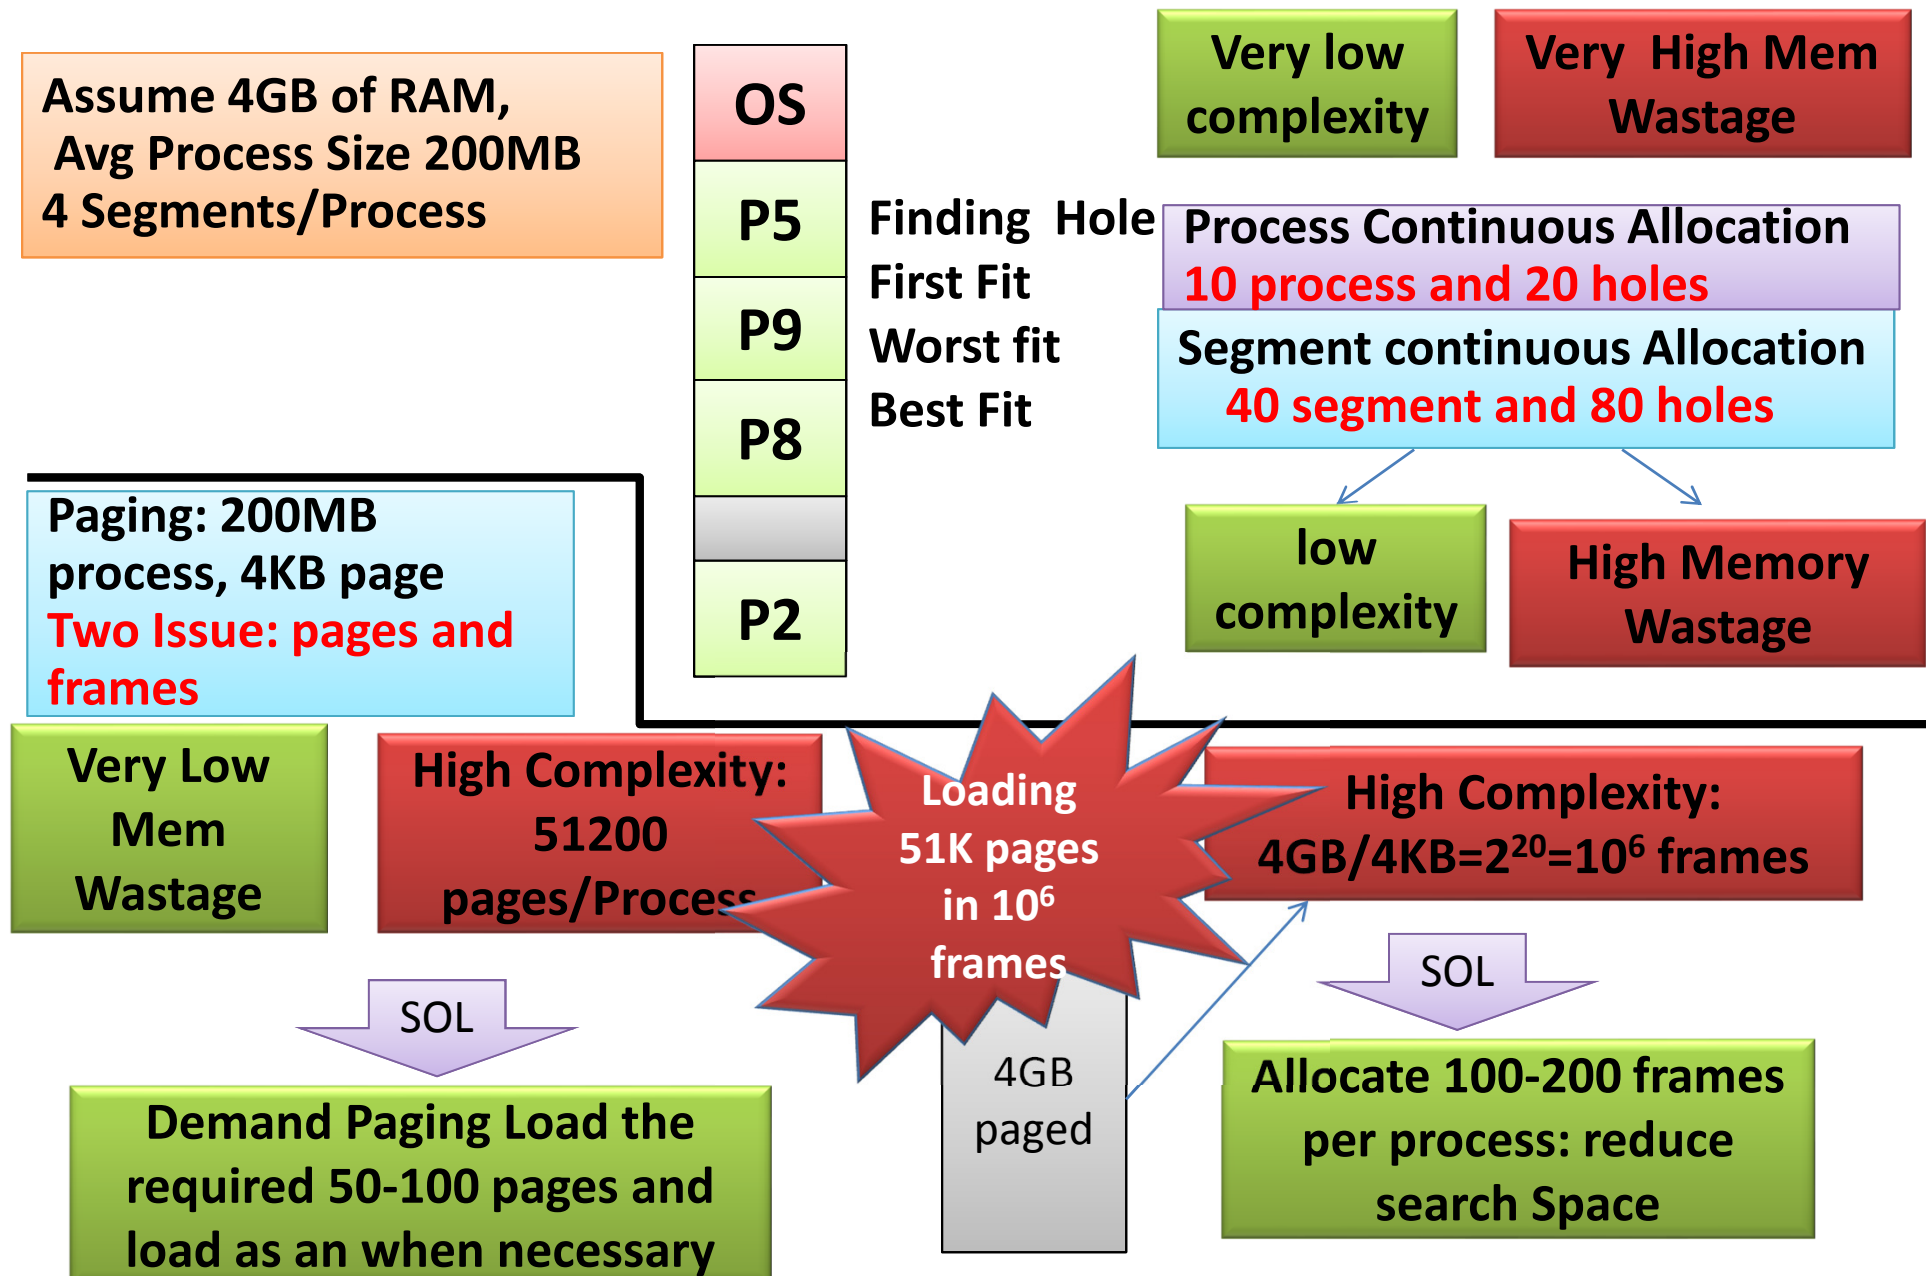
Outline

- Memory Management
 - Continuous Memory allocation
 - Buddy System
 - Segmentation
 - Paging

Protection Issue

- Long term scheduler may put many process in to Ready at a time
 - Where to put ?
 - How to access them ?
 - Is there any efficient way to put and access?
 - How ensure safety and protection?

Memory Allocation: Top Down



Logical vs. Physical Address Space

- Concept of a **logical address space** that is bound to a separate **physical address space**
 - Is central to proper memory management
- **Logical address** – generated by the CPU; also referred to as **virtual address**
- **Physical address** – address seen by the memory unit

Logical vs. Physical Address Space

- **Logical (Virtual) and physical addresses**
 - Are same in compile-time and load-time address-binding schemes
 - Differ in execution-time address-binding scheme
- **Logical address space**
 - Set of all logical addresses generated by a program
- **Physical address space**
 - set of all physical addresses generated by a program

Swapping

- Swapping
 - A process can be **swapped** temporarily out of memory to a backing store
 - And then brought back into memory for continued execution
- Total physical memory space of processes can exceed physical memory
- **Backing store (Ex : HDD/SDD)**
 - fast disk large enough to accommodate copies of all memory images for all users;
 - must provide direct access to these memory images


Swapping

- **Roll out, roll in**
 - Swapping variant used for priority-based scheduling algorithms;
 - Lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time
 - Total transfer time is directly proportional to the amount of memory swapped
- System maintains **queue**
 - Ready-to-run processes which have memory images on disk

Swapping (Cont.)

- Does the swapped out process need to swap back in to same physical addresses?
- Depends on address binding method
 - Plus consider pending I/O to / from process memory space

Schematic View of Swapping



Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory
 - Need to swap out a process and swap in target process
 - Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time 2000 ms + Swap in of same sized process = 4 seconds
 - Swapping component of Context Switch time is 4 seconds

Context Switch Time including Swapping

- Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - Read only memory Discard
 - Code and Declared Read only Data, Constant Area
 - System calls to inform OS of memory use via `request_memory()` and `release_memory()`

Real System :Memory SwapIn/SwapOut

- Server : Utilization is high, always get a crunch.....Supposed to utilized all resources
- PC (4GB RAM, 500GB HDD)
 - Most of the time RAM utilization is bellow 40%
 - Most of the time we get a space for our process to Run..... :) :)
- Mobile (256MB RAM, 8GB SSD)
 - SSD is faster, Read takes less time but write take time, Number write to SSD is limited by a number in SSD life time

Real System :Memory SwapIn/SwapOut

- In PC
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if more than threshold amount of memory allocated
 - Disabled again once memory demand reduced below threshold

Swapping on Mobile Systems

- Not typically supported : Flash memory based
 - Small amount of space
 - Limited number of write cycles (NVRAM, STT RAM, PCM RAM)
 - Poor throughput between flash memory and CPU on mobile platform

Swapping on Mobile Systems

- Instead use other methods to free memory if low
 - iOS *asks* apps to voluntarily relinquish allocated memory
 - Read-only data thrown out and reloaded from flash if needed
 - Failure to free can result in termination
 - Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
 - Both OSes support paging (will be discussed)


Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory


Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing OS code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*

Base and Limit Registers



Hardware Address Protection



Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Fixed size Partition (Partition Size (PS))**
 - Num process to support is limited if PS is bigger
 - Supporting a big size process is limited if PS is smaller
- **Variable-partition** sizes for efficiency
 - Sized to a given process' needs
- **Hole** – block of available memory
 - Holes of various size are scattered throughout memory

Buddy System

- A compromise between fixed size and variable size allocation
- Memory is allocated that are a power of 2
- Initially a single allocation unit
- A process is allocated a unit memory whose size is the smallest power of 2 larger than the size of process
 - 50K process would be place in 64K space
- If no allocation unit exist of that size, the smallest available allocation larger than the process will be split into two **“Buddy unit”** of $\frac{1}{2}$ size

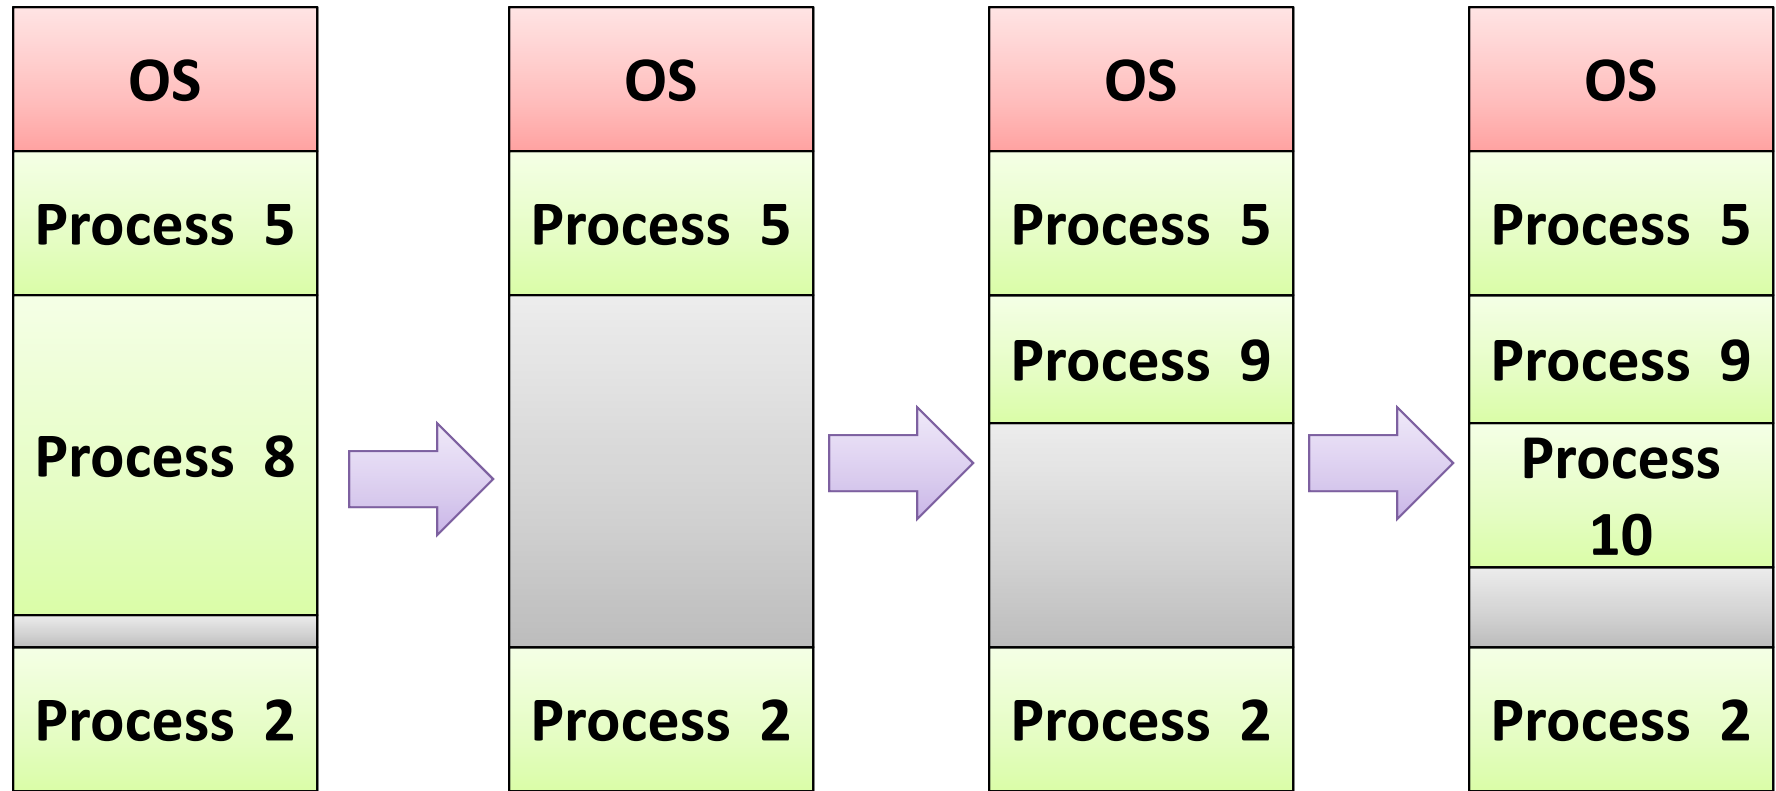
Example

Action	Memory					
Start	1024K					
Arrive A (150K)	A	256K		512K		
Arrive B (100K)	A	B	128K	512K		
Arrive C (50K)	A	B	C	64 k	512K	
Release B	A	128K	C	64 k	512K	
Arrive D (200K)	A	128K	C	64 k	D	256K
Arrive E (60K)	A	128K	C	E	D	256K
Release C	A	128K	64 k	E	D	256K
Release A	256K	128K	64 k	E	D	256K
Release E	512K				D	256K
Release D	1024K					

Multiple-partition allocation

- When a process arrives
 - OS allocate memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- OS maintains information about
 - a) allocated partitions
 - b) free partitions (hole)

Multiple-partition allocation



Three
Processes
Running

Process 8
Finished

Process 9
Arrives

Process 10
Arrives

Dynamic Storage-Allocation Problem

- How to satisfy a request of size n from a list of free holes?
 - **First-fit**: Allocate the *first* hole that is big enough
 - **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, **unless ordered by size**
 - Produces the smallest leftover hole
 - **Worst-fit**: Allocate the *largest* hole; must also search entire list
 - Produces the largest leftover hole
- First-fit and best-fit better than worst-fit in terms of
 - speed (FF is better as compared to WF) and
 - storage utilization (BF is better as compared to WF)

Fragmentation

- **External Fragmentation** – Total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory
 - This size difference is memory internal to a partition, but not being used

Fragmentation: First Fit

- Analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable
 - **50-percent rule**
 - Once memory is almost full
 - Events (processes arrival and departure) are alternate

Fragmentation (Cont.)


- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - I/O problem
 - Latch job in memory while it is involved in I/O
 - Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems

Segmentation

Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
 - A segment is a logical unit such as:
main program,
procedure/function/method,
object, local variables, global variables
common block, stack, symbol table
arrays

User's View of a Program




Logical address


Segmentation : Key Rule

- Instead of allocation memory for whole process
- Divide the program into smaller block call segments
 - User view : Already segmented
 - Finer granularity, less fragmentation
- **Mustard in Bag Vs Brinjal in Bag**
- Each of which is allocated to memory independently
- Segments are variable size

Logical View of Segmentation



user space



physical memory space

Segmentation Architecture

- Logical address consists of a two tuple:
 <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segments reside in memory
 - **limit** – specifies the length of the segment

Segmentation Architecture

- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;

segment number **s** is legal if **$s < \text{STLR}$**

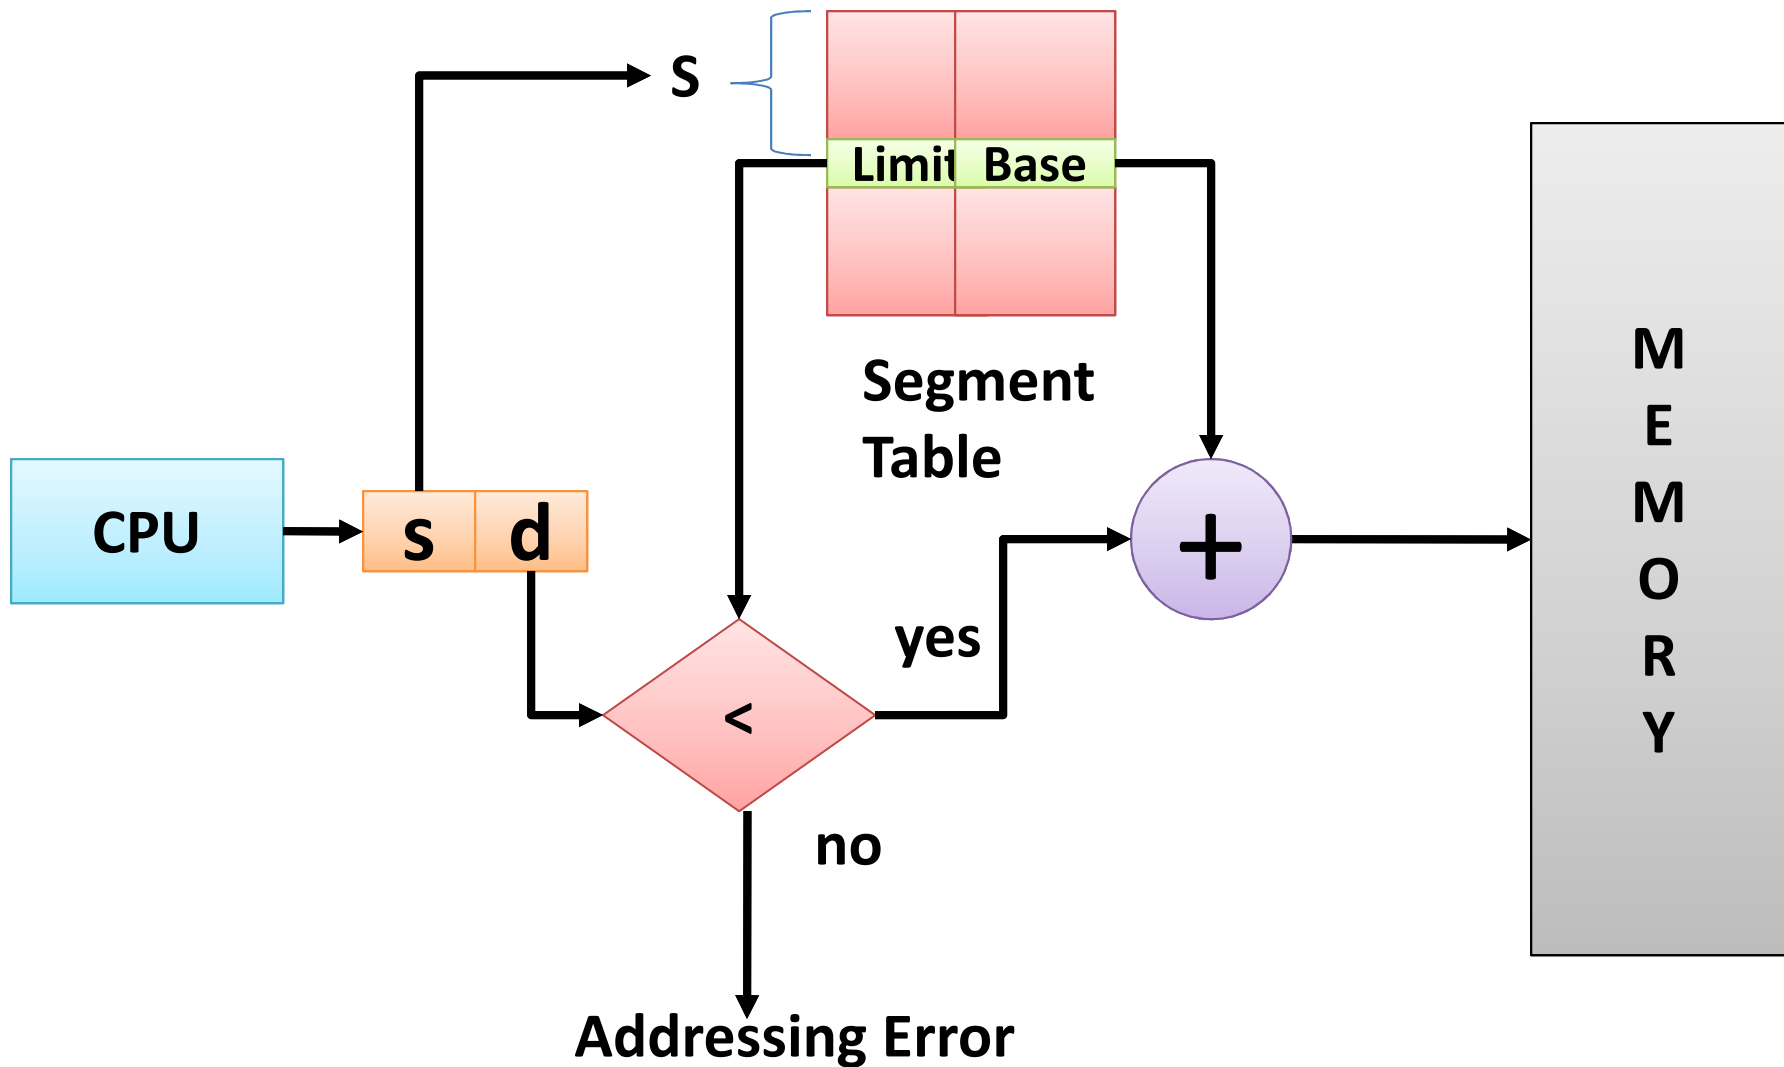
Every memory access get
translated to Two accesses



Segmentation Architecture (Cont.)

- Protection
 - With each entry in segment table associate:
 - validation bit = 0 \Rightarrow illegal segment
 - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem

Segmentation Hardware



Segmentation and Paging

- Divide the program into smaller block call segments : User view, already segmented
 - Finer granularity, less fragmentation
 - **Mustard in Bag Vs Brinjal in Bag**
- Divide the program in to smaller but uniform size unit called page
 - A program may contain many pages
 - Last page of program may be partially filled
- **Divide the memory in to smaller size units called Frame**
- Page get mapped to Frame

Thanks