

**CS343: Operating System**

**Memory Management:  
Segmentation and Paging**

**Lect31 : 27th Oct 2023**

**Dr. A. Sahu**


**Dept of Comp. Sc. & Engg.**

**Indian Institute of Technology Guwahati**

# Outline

- Memory Management
  - Continuous Memory allocation
  - Buddy System
  - Segmentation
  - Paging

# Memory Allocation: Top Down



# Fragmentation

- **External Fragmentation** – Total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory
  - This size difference is memory internal to a partition, but not being used

# Fragmentation: First Fit

- Analysis reveals that given  $N$  blocks allocated,  $0.5 N$  blocks lost to fragmentation
  - $1/3$  may be unusable
  - **50-percent rule**

# Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible *only* if relocation is dynamic, and is done at execution time
  - I/O problem
    - Latch job in memory while it is involved in I/O
    - Do I/O only into OS buffers
- Now consider that backing store has same fragmentation problems


# Segmentation

# Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
  - A segment is a logical unit such as:  
main program,  
procedure/function/method,  
object, local variables, global variables  
common block, stack, symbol table  
arrays



# User's View of a Program




**Logical address**


# Segmentation : Key Rule

- Instead of allocation memory for whole process
- Divide the program into smaller block call segments
  - User view : Already segmented
  - Finer granularity, less fragmentation
- **Mustard in Bag Vs Brinjal in Bag**
- Each of which is allocated to memory independently
- Segments are variable size

# Logical View of Segmentation



**user space**



**physical memory space**

# Segmentation Architecture

- Logical address consists of a two tuple:  
    <segment-number, offset>,
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – contains the starting physical address where the segments reside in memory
  - **limit** – specifies the length of the segment

# Segmentation Architecture

- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;

segment number  **$s$**  is legal if  **$s < \text{STLR}$**

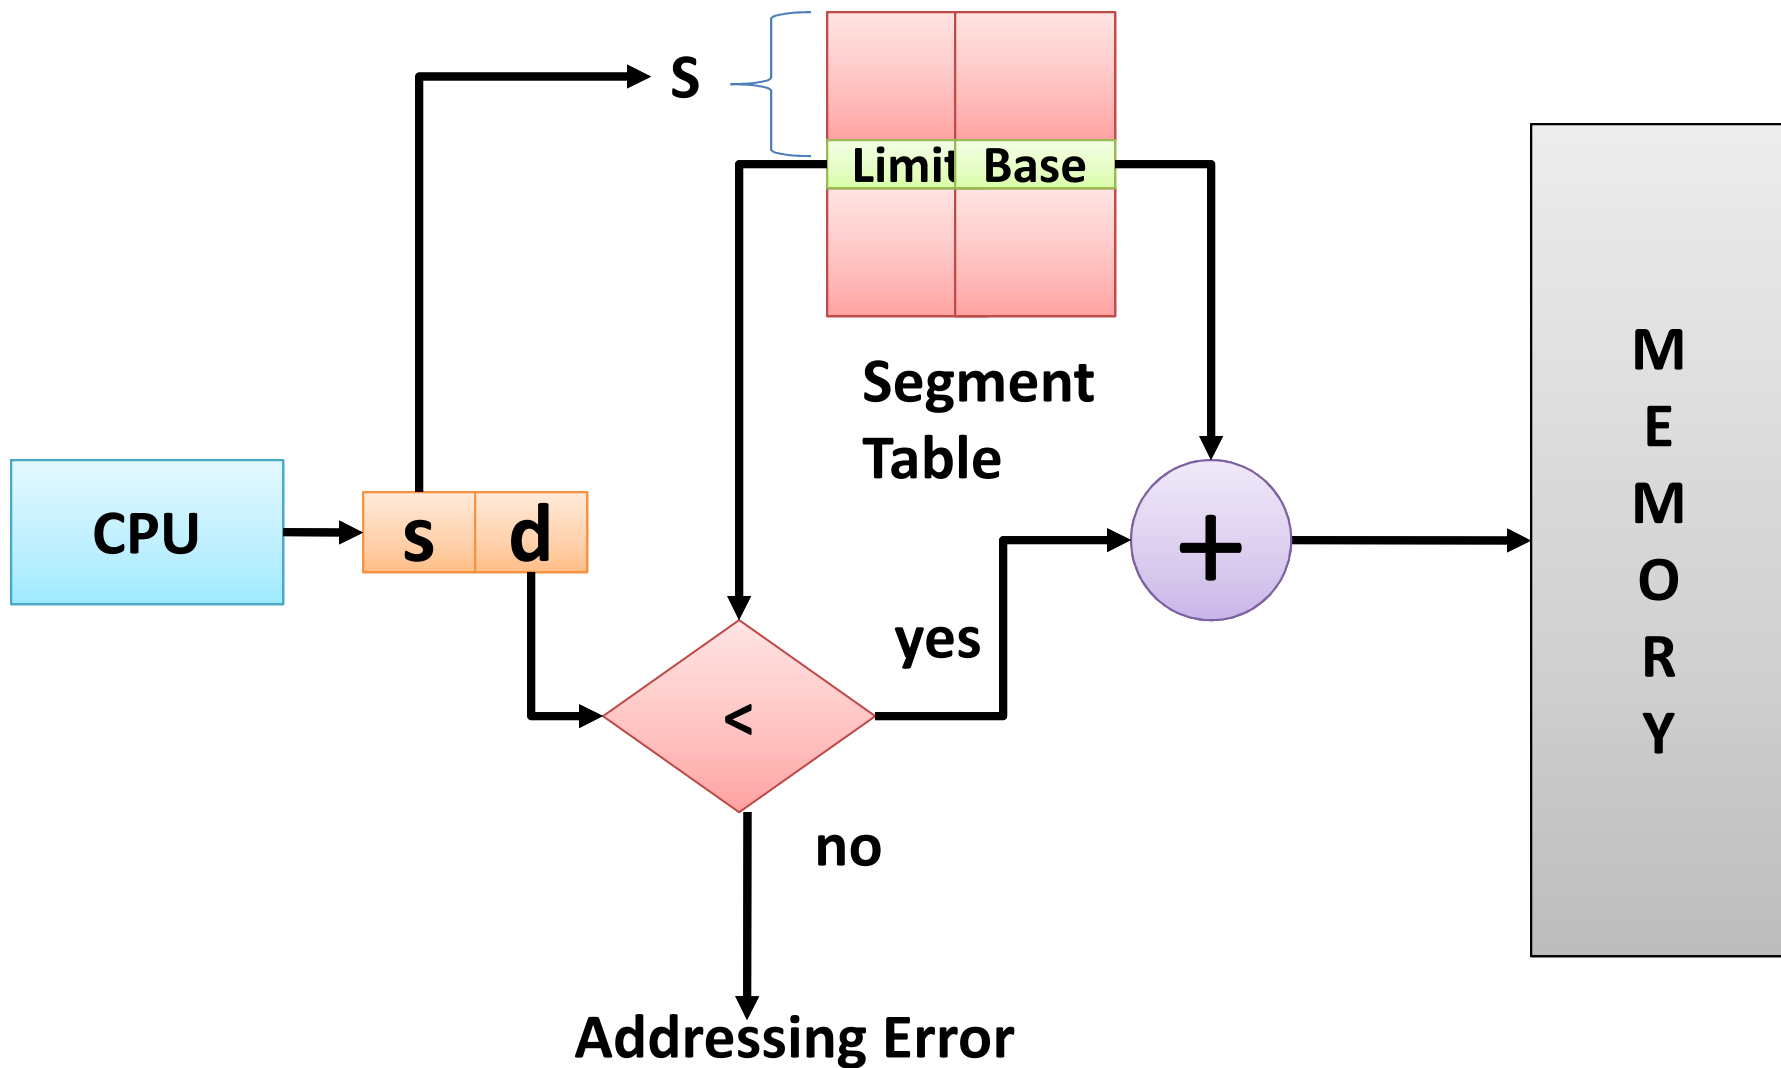
Every memory access get  
translated to Two accesses



# Segmentation Architecture (Cont.)

- Protection
  - With each entry in segment table associate:
    - validation bit = 0  $\Rightarrow$  illegal segment
    - read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram

# Segmentation Hardware



# Segmentation and Paging

- Divide the program into smaller block call segments : User view, already segmented
  - Finer granularity, less fragmentation
    - **Mustard in Bag Vs Brinjal in Bag**
- Divide the program in to smaller but uniform size unit called page
  - A program may contain many pages
  - Last page of program may be partially filled
- Divide the memory in to smaller size units called Frame
- Page get mapped to Frame



# Paging

# Paging


- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
  - Size is power of 2, between 512 bytes and 16 Mbytes , **getpagesize() in C → Demo**

# Paging

- Divide logical/program memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size  **$N$**  pages, need to find  **$N$**  free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation


# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit




- $m-n$                        $n$
- For given logical address space  $2^m$  and page size  $2^n$

# Paging Hardware



# Paging Model of Logical and Physical Memory



# Paging (Cont.)


- Calculating internal fragmentation
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
  - 35 pages + 1,086 bytes
  - Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
  - Worst case fragmentation = 1 frame – 1 byte
- Another Example: Suppose you want to withdraw Rs 533 from ATM
  - ATM allow only multiple of 100
  - Either 500 or 600, Fragmentation  $100 - 33 = 67$
  - **Worst case, need to withdraw Rs 501, Frag.=  $100 - 1 = 99$**

# Paging (Cont.)

- Calculating internal fragmentation Ctd..
  - On average fragmentation =  $1 / 2$  frame size
  - So small frame sizes desirable?
  - But each page table entry takes memory to track
  - Page sizes growing over time
    - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory



# Page Table Example



Process B


0	3
1	7

Page Table

page number	page offset
p	d

$m-n=3$

$n=4$




Process A

0	1
1	4

Page Table

Free list  
0, 2, 5, 7

$b=7$





Physical  
Memory

# Free Frames

free-frame list

14  
13  
18  
20  
15




(a)

Before allocation

free-frame list

15



0	14
1	13
2	18
3	20

new-process page table




(b)

After allocation

# Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
  - One for the page table and one for the data / instruction

# Implementation of Page Table



- Page Table Length
- Two memory accesses per data/inst access.
  - Solution? *Associative Registers*

# Implementation of Page Table (Cont.)

- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
  - Get implemented in Processor as CAM
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch

# Implementation of Page Table (Cont.)

- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access

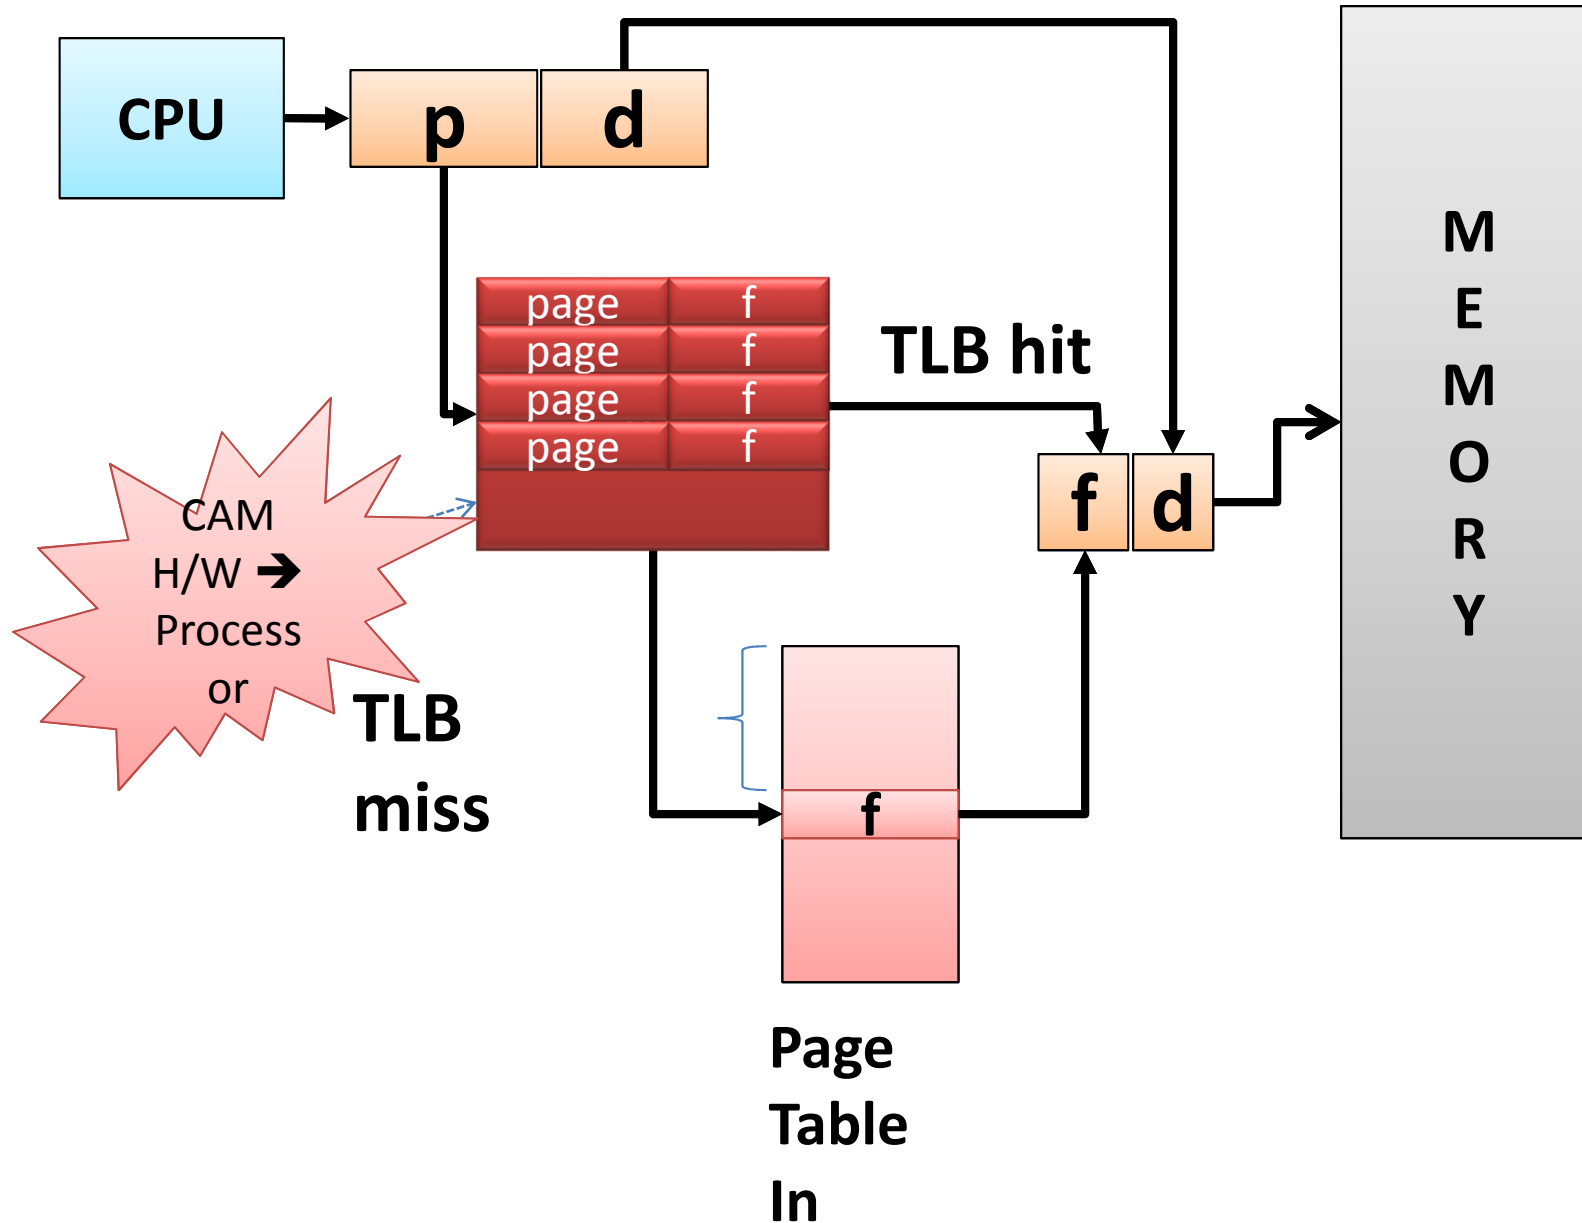
# Associative Memory: CAM

- Associative memory – parallel search

Page #	Frame #
1	2
3	15
2	12

- **Content Addressable memory : CAM**
- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory

# Paging Hardware With TLB





# Thanks