

CS343: Operating System

Memory Management

Lect32 : 30th Oct 2023

Dr. A. Sahu


Dept of Comp. Sc. & Engg.

Indian Institute of Technology Guwahati

Outline

- Memory Management
 - Paging
 - Virtual memory
 - Demand Paging
 - Page Replacement
 - Frame Allocation

Memory Allocation: Top Down



Paging

Paging


- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes , **getpagesize() in C → Demo**

Paging

- Divide logical/program memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size **N** pages, need to find **N** free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit

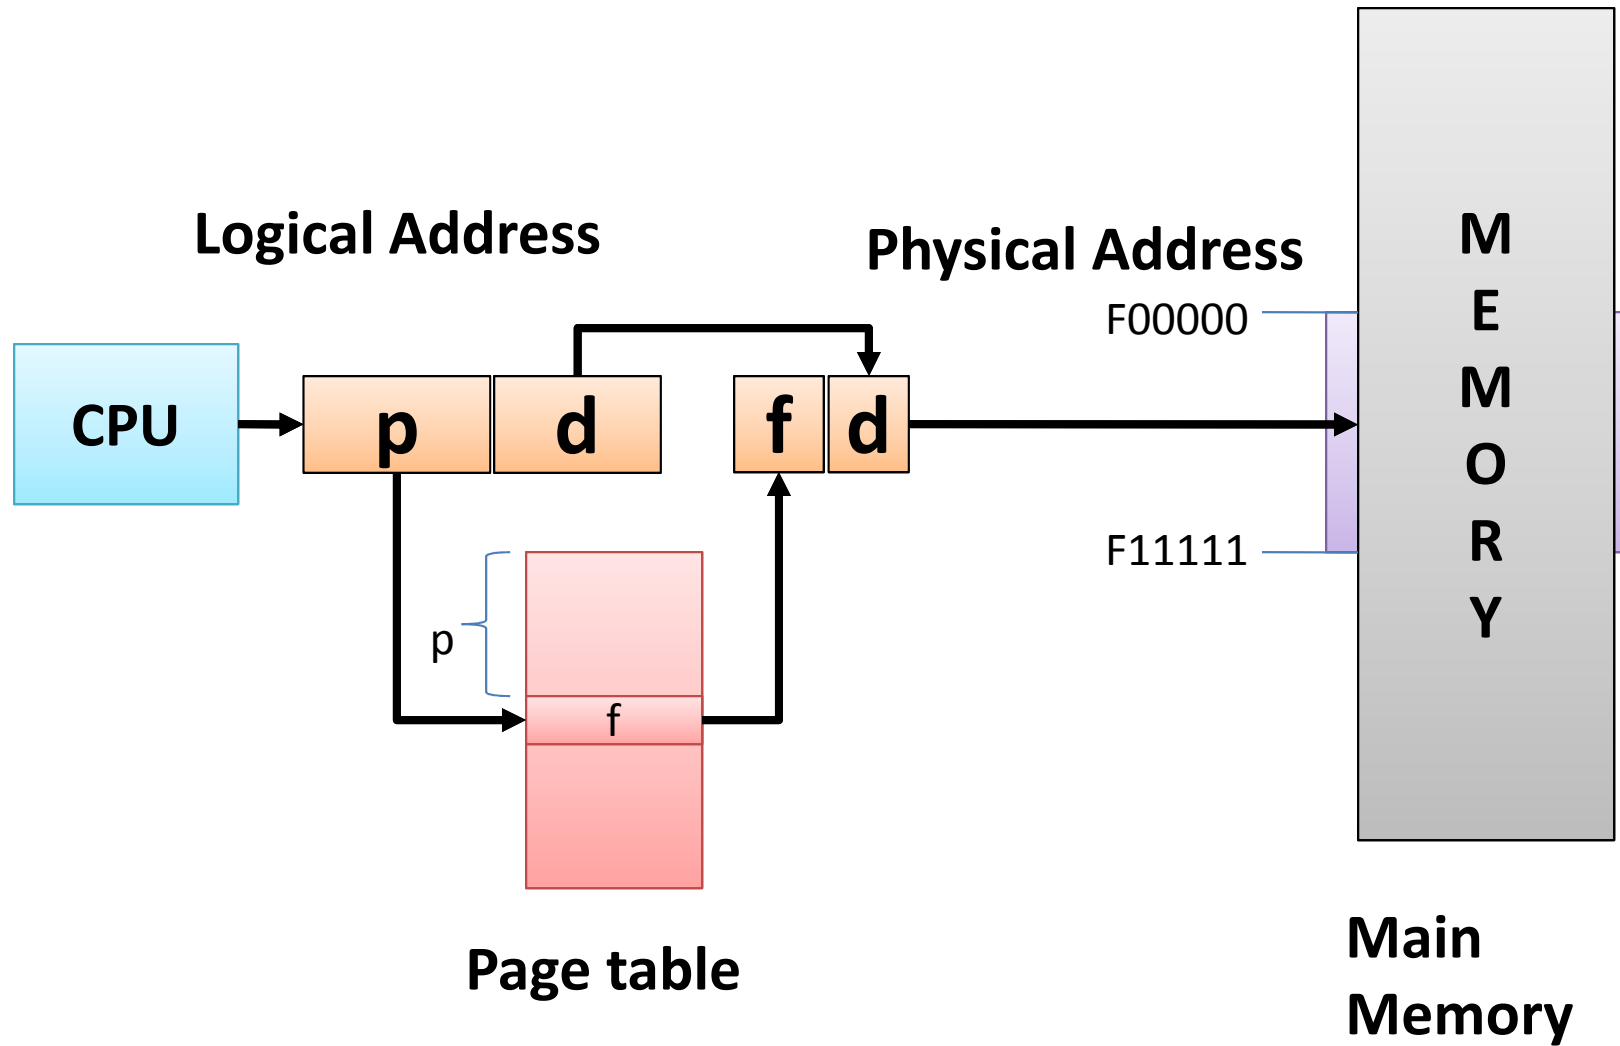


$m-n$

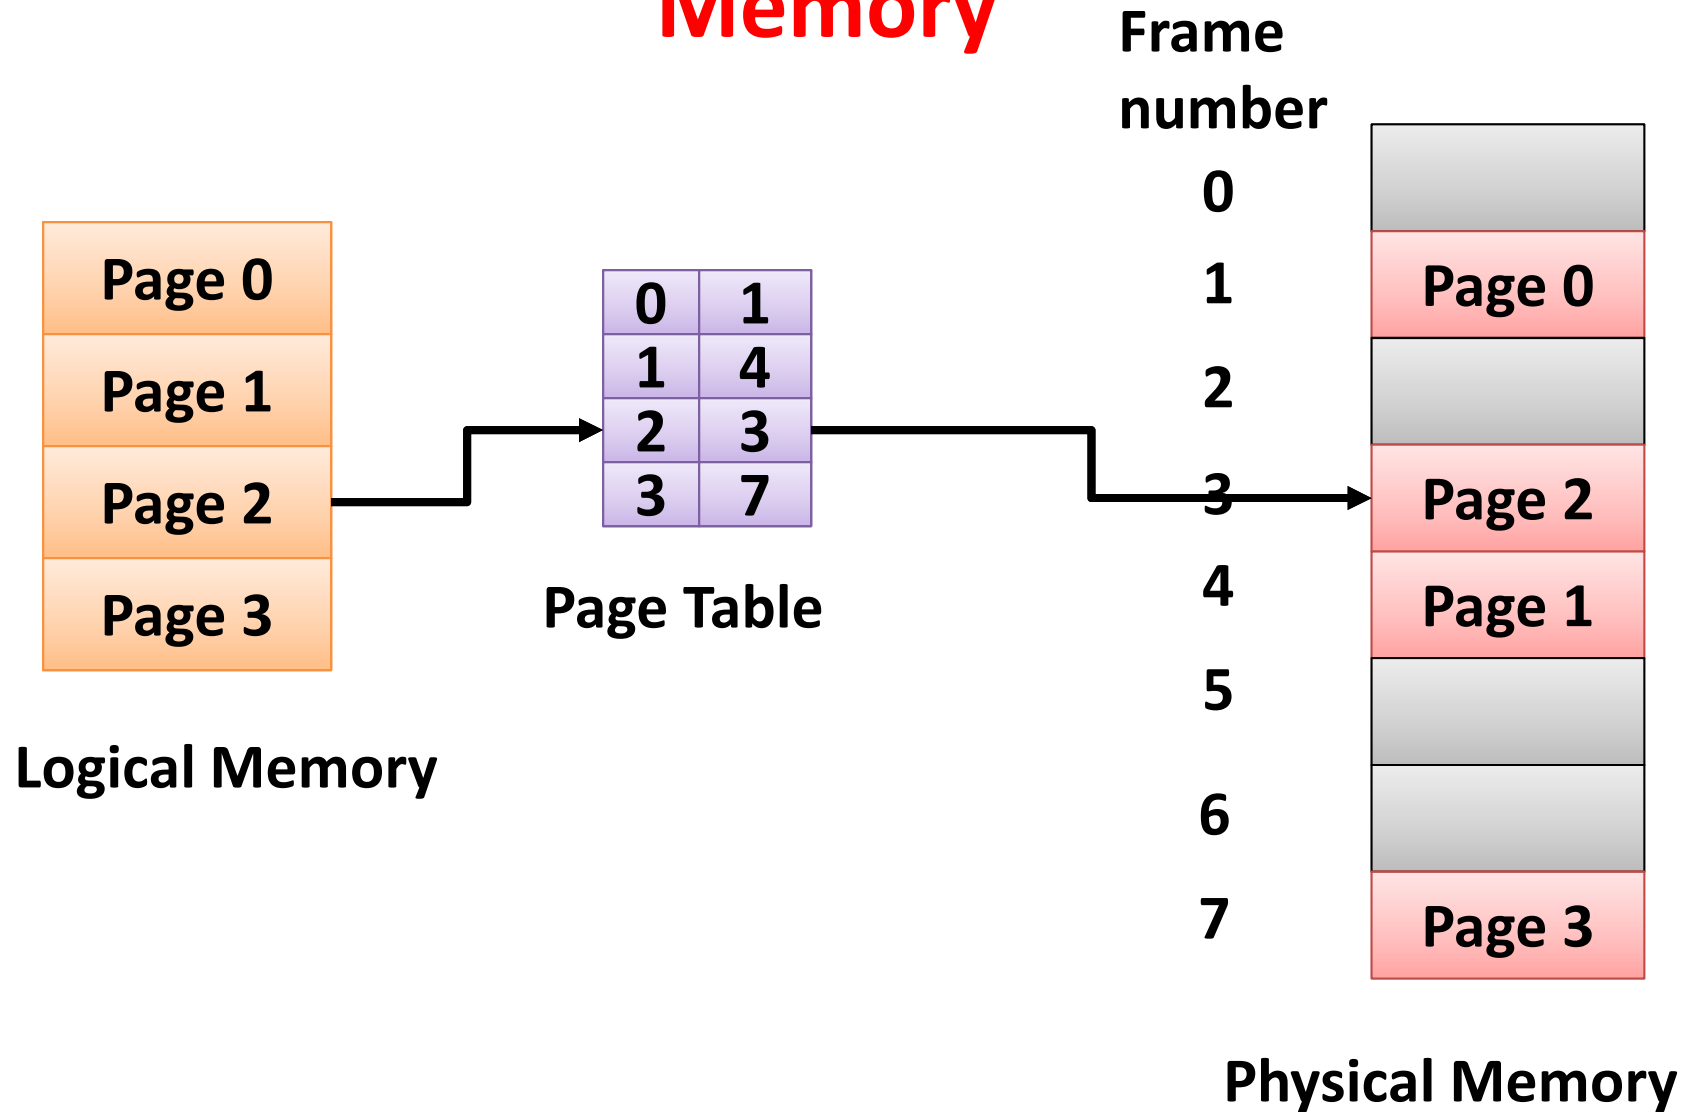
n

- For given logical address space 2^m and page size 2^n

Paging Hardware



Paging Model of Logical and Physical Memory




Paging (Cont.)

- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte

Paging (Cont.)

- Calculating internal fragmentation Ctd..
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Page Table Example



Process B


0	3
1	7

Page Table

page number	page offset
p	d

$$m-n=3$$

$$n=4$$




Process A

0	1
1	4

Page Table

Free list
0, 2, 5, 7

$b=7$




Physical
Memory

Free Frames


free-frame list

14
13
18
20
15




(a)

Before allocation



free-frame list

15



(b)

new-process page table

0	14
1	13
2	18
3	20




After allocation

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction

Implementation of Page Table



- Page Table Length
- Two memory accesses per data/inst access.
 - Solution? *Associative Registers/CAM*

Implementation of Page Table (Cont.)

- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch

Implementation of Page Table (Cont.)

- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access

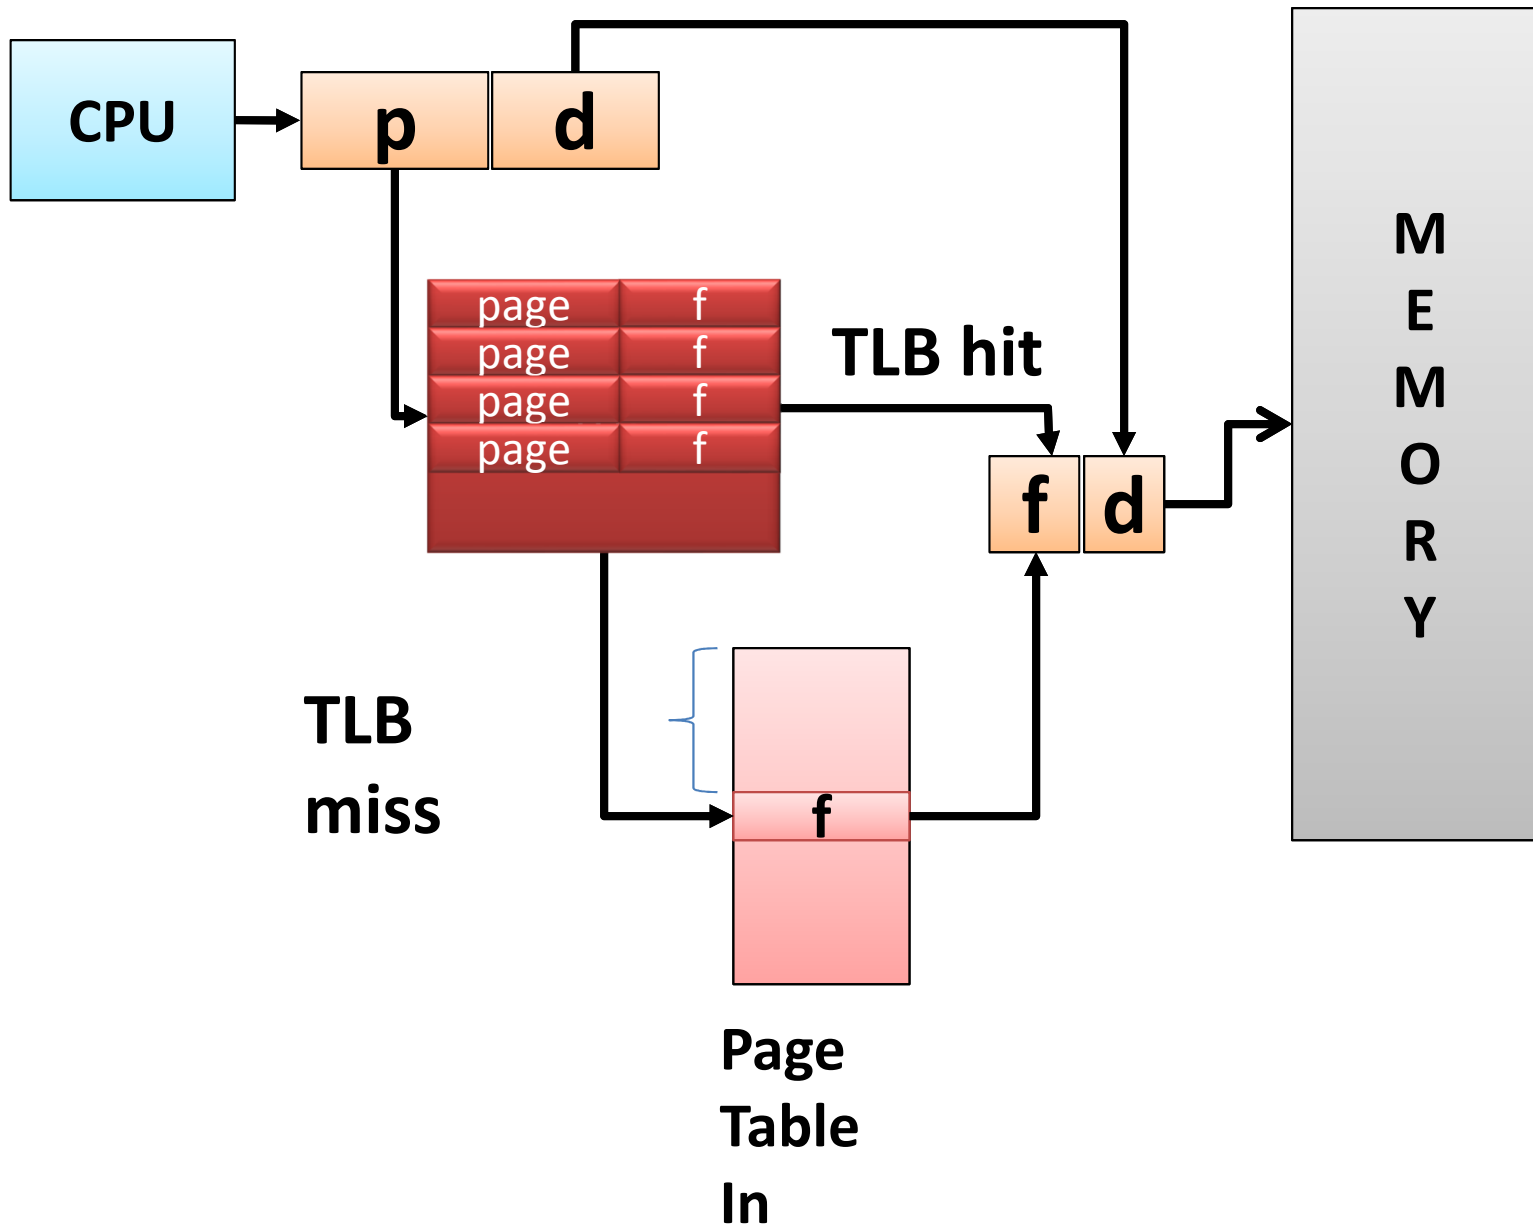
Associative Memory

- Associative memory – parallel search

Page #	Frame #
1	2
3	15
2	12

- CAM : Content Addressable Memory
- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



Effective Access Time

- Associative Lookup = ϵ time unit
 - Can be $< 10\%$ of memory access time
- Hit ratio = h
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider $h = 80\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- **Effective Access Time (EAT)**

$$\begin{aligned} \text{EAT} &= (1 + \epsilon) h + (2 + \epsilon)(1 - h) \\ &= 2 + \epsilon - h \end{aligned}$$

Effective Access Time

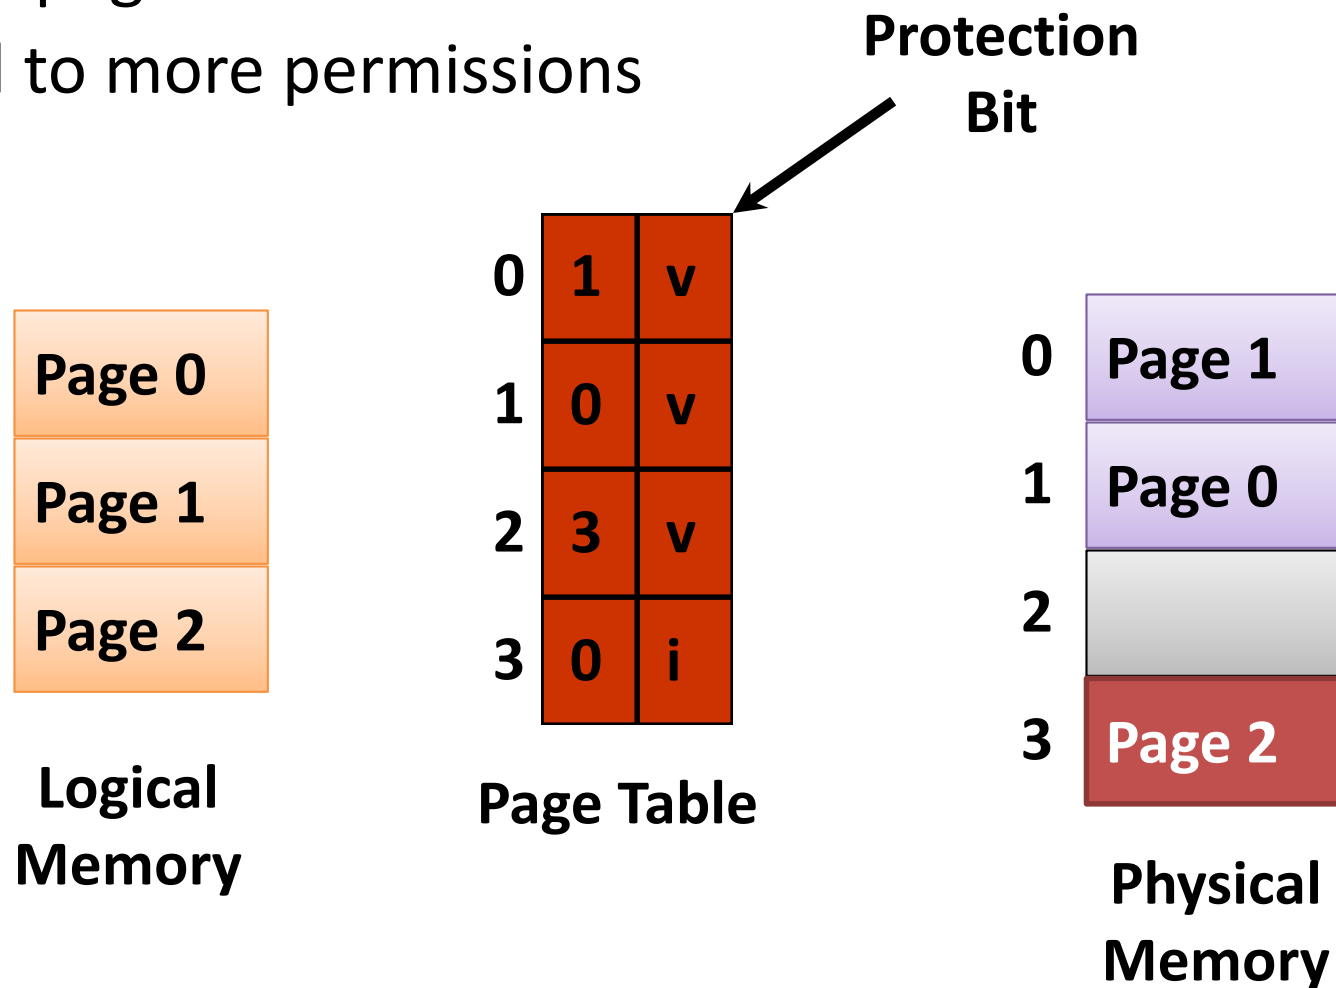
- Consider $h = 80\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.80 \times (100+20) + 0.20 \times (200+20) = 140\text{ns}$
- Consider more realistic hit ratio $\rightarrow h = 99\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.99 \times (120) + 0.01 \times 220 = 121\text{ns}$

Memory Protection

- Associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space: a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel

Protection





- Protection bits with each frame
- Store in page table
- Expand to more permissions



Shared Pages

- **Shared code**
 - One copy of read-only (**reentrant**) code shared among processes
 - (i.e., text editors, compilers, window systems)
 - Also useful for interprocess communication if sharing of read-write pages is allowed
- **Private code and data**
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example



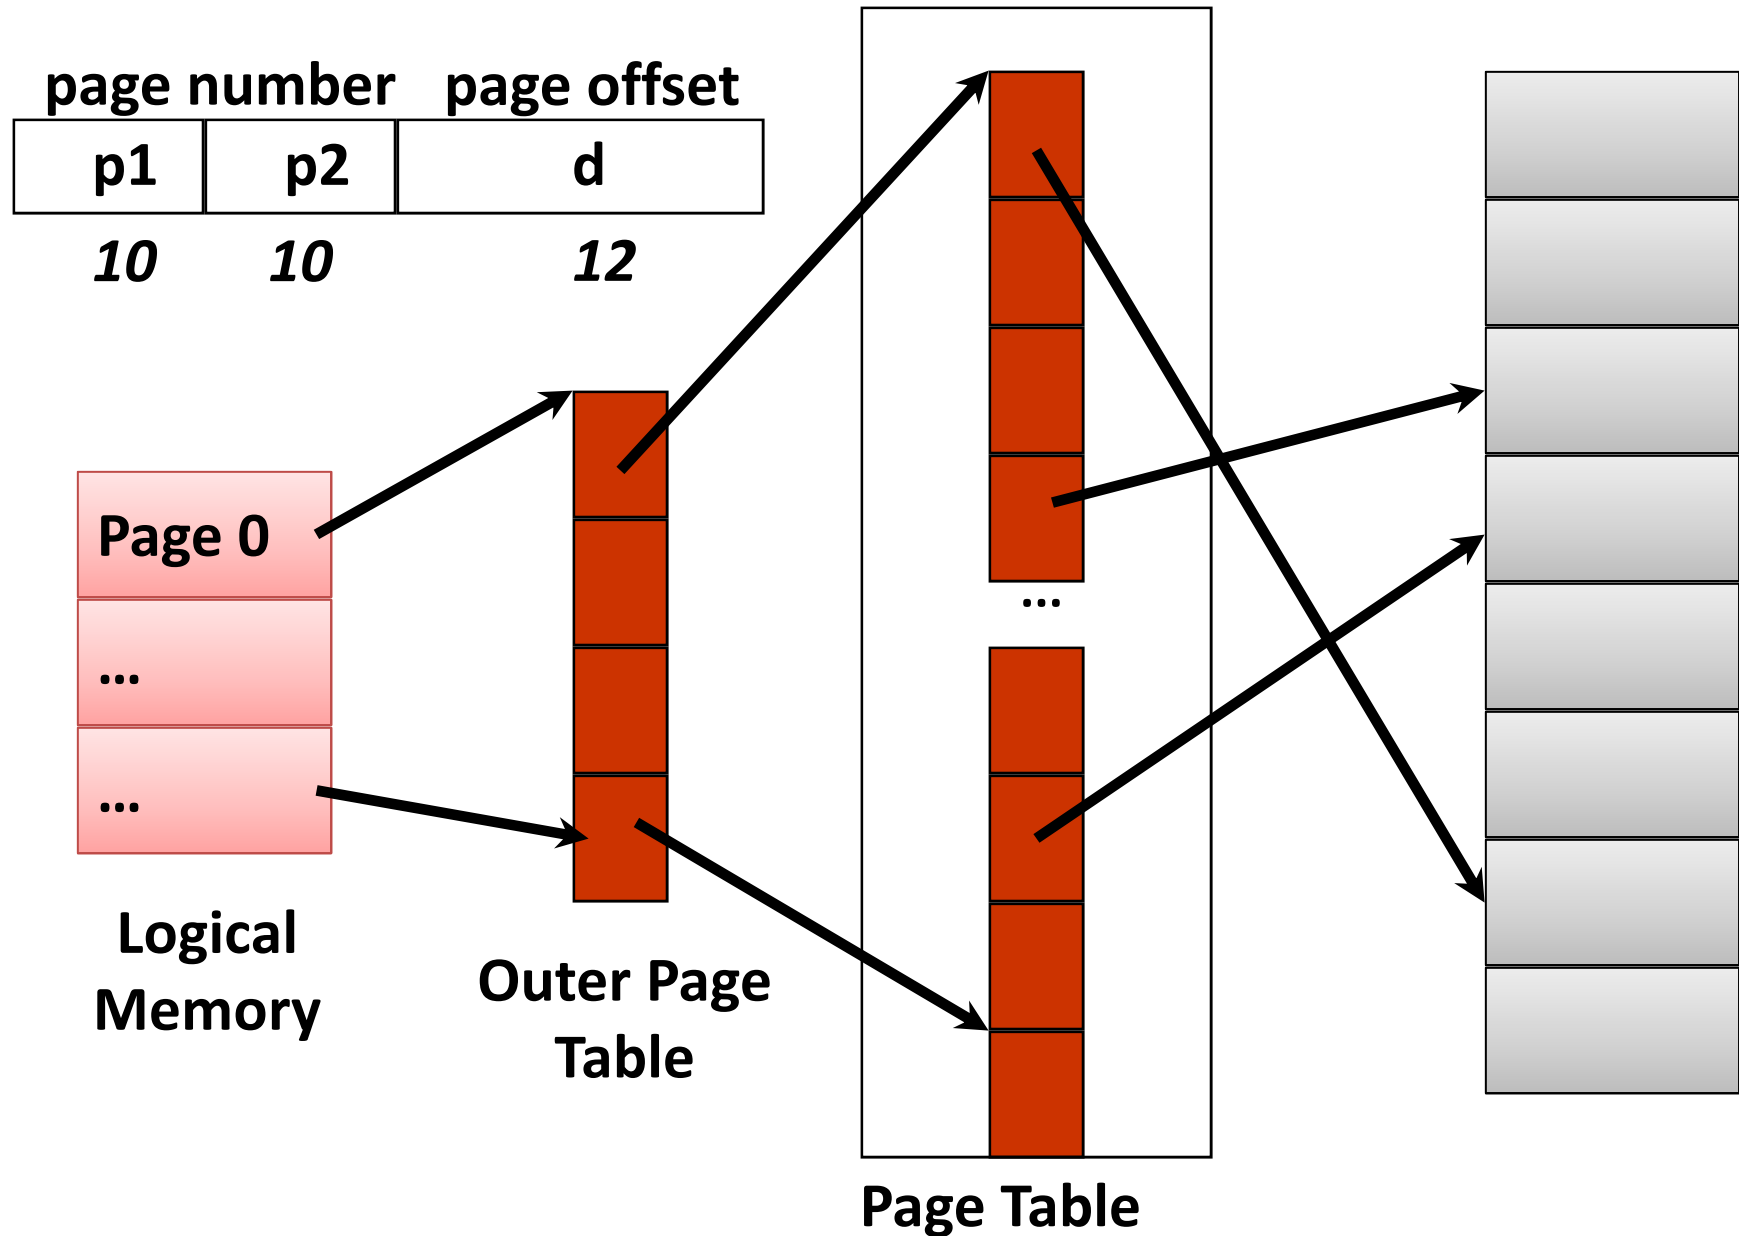
Large Address Spaces

- Typical logical address spaces:
 - 4 Gbytes $\Rightarrow 2^{32}$ address bits (4-byte address)
- Typical page size:
 - 4 Kbytes = 2^{12} bits
- Page table may have:
 - $2^{32} / 2^{12} = 2^{20} = 1\text{million}$ entries
- Each entry 4 bytes $\Rightarrow 4\text{MB}$ per process!
Linux command “`$ps -A |wc -l`” show 400 process running most of the time
- **Do not want that all in RAM**
 - **$400 * 4\text{MB} = 1.6\text{GB}$ of RAM will be occupied by Page Tables**
- **Solution?**
 - **Page the Page table (Multilevel paging), Hashing, and Inverted Page Table**

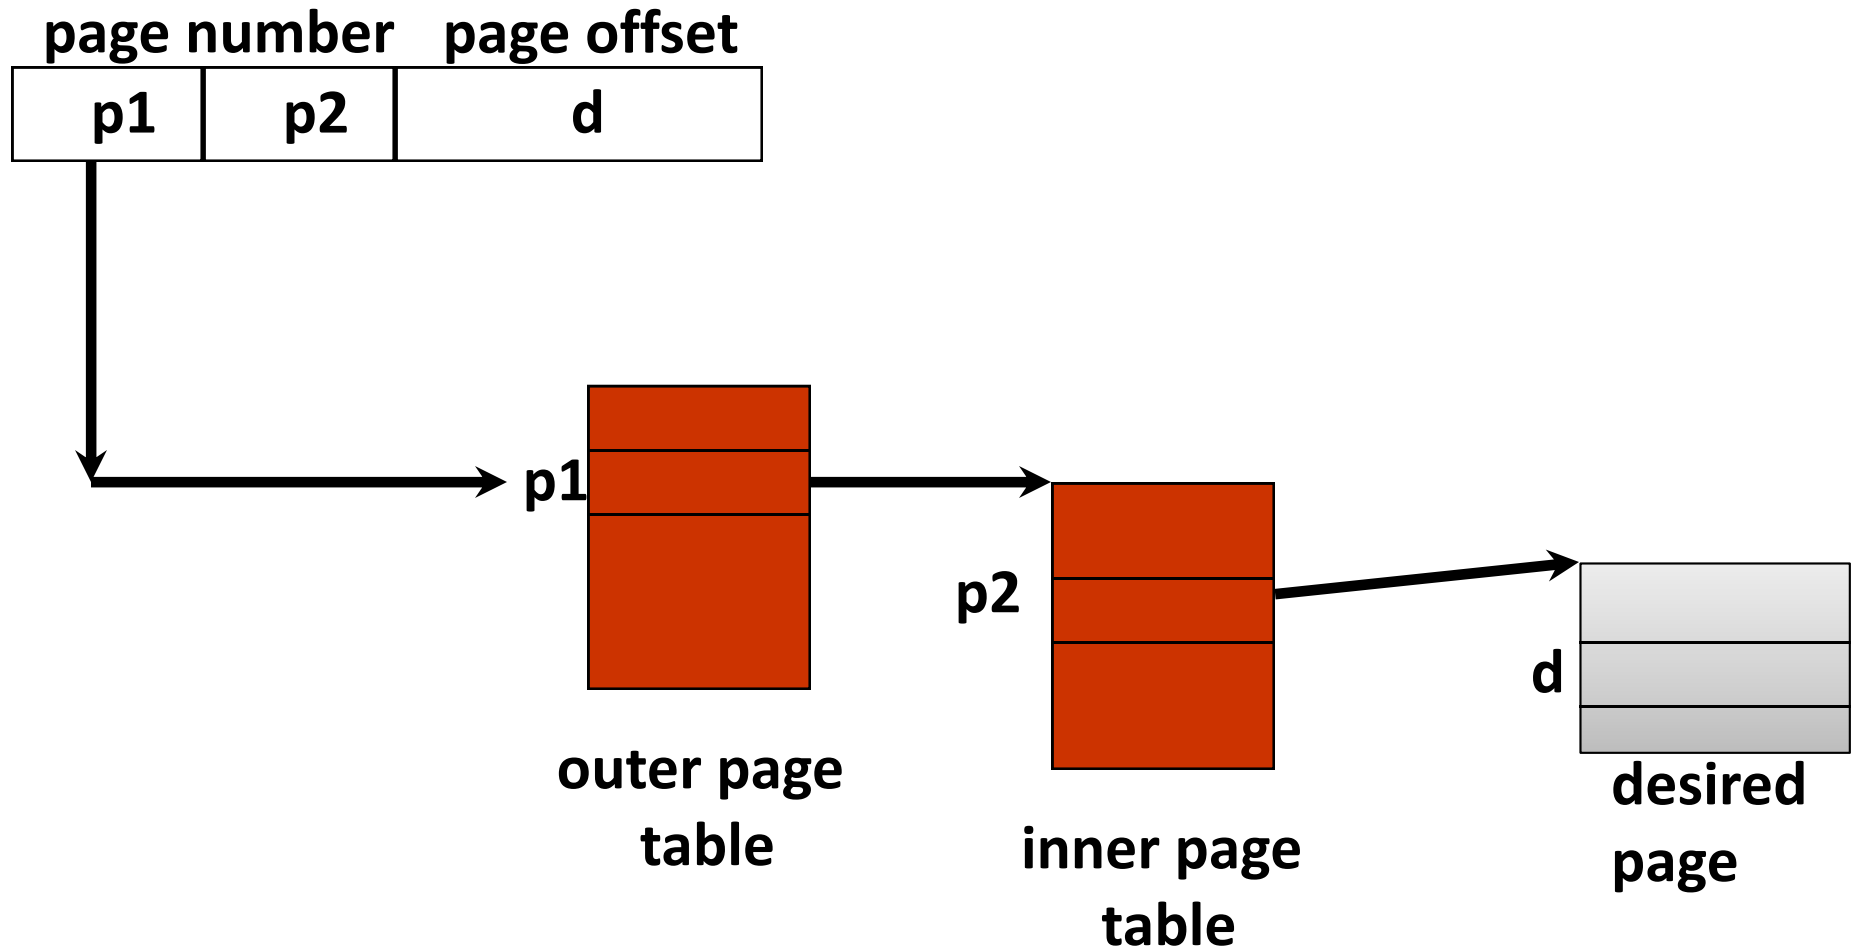
Large Address Spaces

- Do not want that all 4MB in RAM
- Solution?
 - Page the Page table (Multilevel paging)
 - We would not want to allocate PT contiguously on memory, divide the PT in to smaller pieces (ML paging)
 - Hashing: if some pages are not to be used of a process is beneficial
 - Inverted Page Table

Multilevel Paging



Multilevel Paging Translation




Hashed Page Tables

- Common in address spaces > 32 bits
 - **Modern Ubuntu and Windows uses 48bit**
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains
 - The virtual page number
 - The value of the mapped page frame
 - A pointer to the next element

Hashed Page Tables

- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted
- Variation for 64-bit addresses is **clustered page tables**
 - Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)

Hashed Page Table



Inverted Page Table

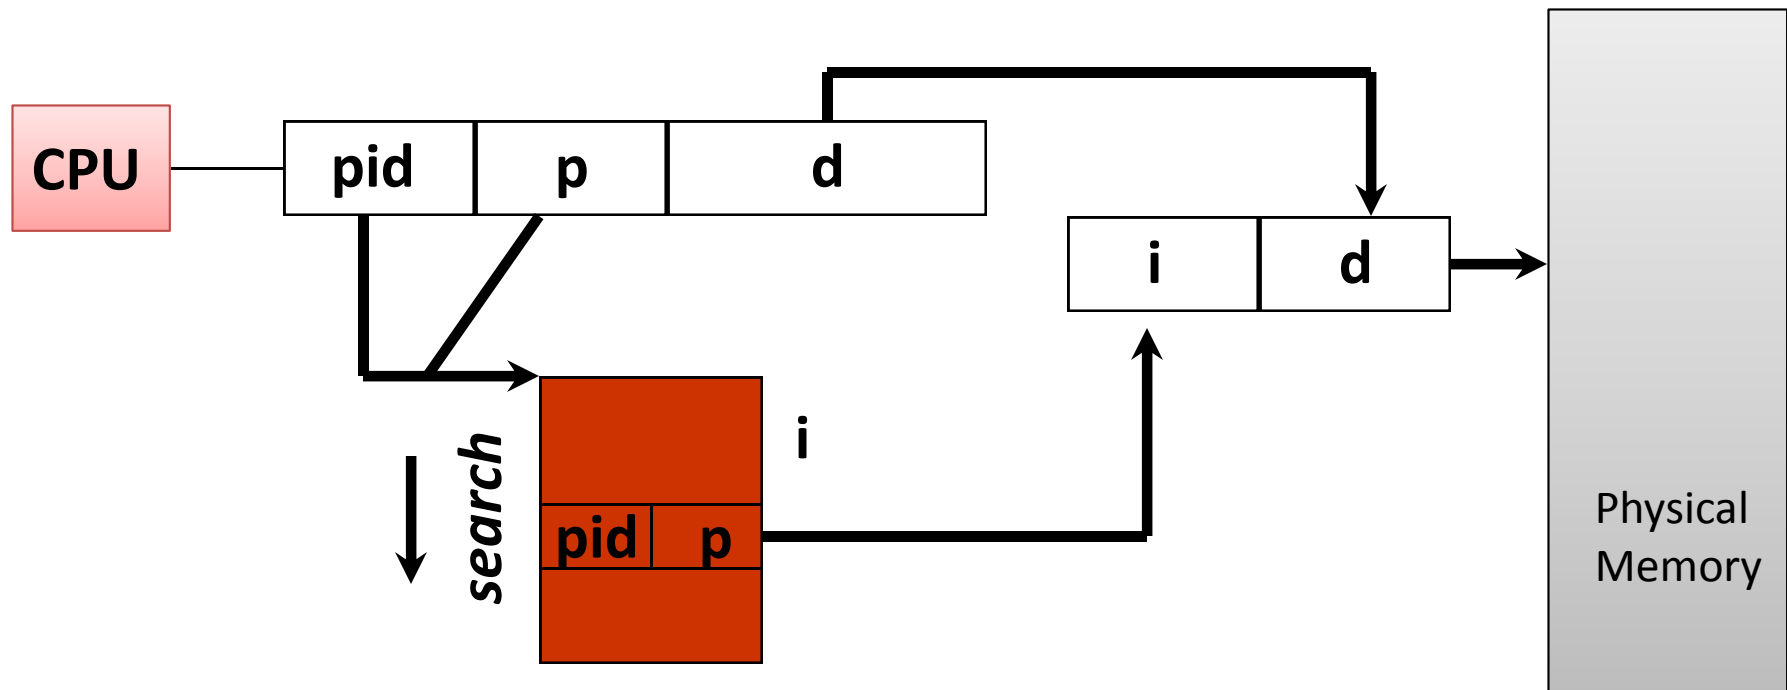
- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page/frame of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page

Inverted Page Table

- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
 - TLB can accelerate access
- But how to implement shared memory?
 - One mapping of a virtual address to the shared physical address

Inverted Page Table

- Page table maps to physical addresses



- Still need page per process --> backing store
- Memory accesses longer! (search + swap)

Thanks