

**CS343: Operating System**

**Virtual Memory, Demand  
Paging, Page Replacement**

**Lect34 : 2<sup>nd</sup> Nov 2023**

**Dr. A. Sahu**

**Dept of Comp. Sc. & Engg.**

**Indian Institute of Technology Guwahati**

# Outline

- Memory Management
  - Paging
  - Virtual memory
  - Demand Paging
  - **Page Replacement**
  - **Frame Allocation**

# Memory Allocation: Top Down

Assume 4GB of RAM,  
Avg Process Size 200MB  
4 Segments/Process

Very low  
complexity

Very High Mem  
Wastage

Finding Hole  
First Fit  
Worst fit  
Best Fit

Process Continuous Allocation  
**10 process and 20 holes**

Segment continuous Allocation  
**40 segment and 80 holes**

Paging: 200MB  
process, 4KB page  
**Two Issue: pages and  
frames**

low  
complexity

High Memory  
Wastage

Very Low  
Mem  
Wastage

High Complexity:  
51200  
pages/Process

Loading  
51K pages  
in  $10^6$   
frames

High Complexity:  
 $4\text{GB}/4\text{KB}=2^{20}=10^6$  frames

SOL

SOL

Demand Paging Load the  
required 50-100 pages and  
load as an when necessary

4GB  
paged

Allocate 100-200 frames  
per process: reduce  
search Space

# Demand Paging

# Demand Paging

- Could bring entire process into memory at load time
- Or bring a page into memory only when it is needed
  - Less I/O needed, no unnecessary I/O
  - Less memory needed
  - Faster response
  - More users

# Demand Paging

- Similar to paging system with swapping
- Page is needed  $\Rightarrow$  reference to it
  - invalid reference  $\Rightarrow$  abort
  - not-in-memory  $\Rightarrow$  bring to memory
- **Lazy swapper** – never swaps a page into memory unless page will be needed
  - Swapper that deals with pages is a **pager**

# Page Table : Some Pages Are Not in Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


Logical  
Memory

#P	#F	V
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

Page Table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	

Main memory

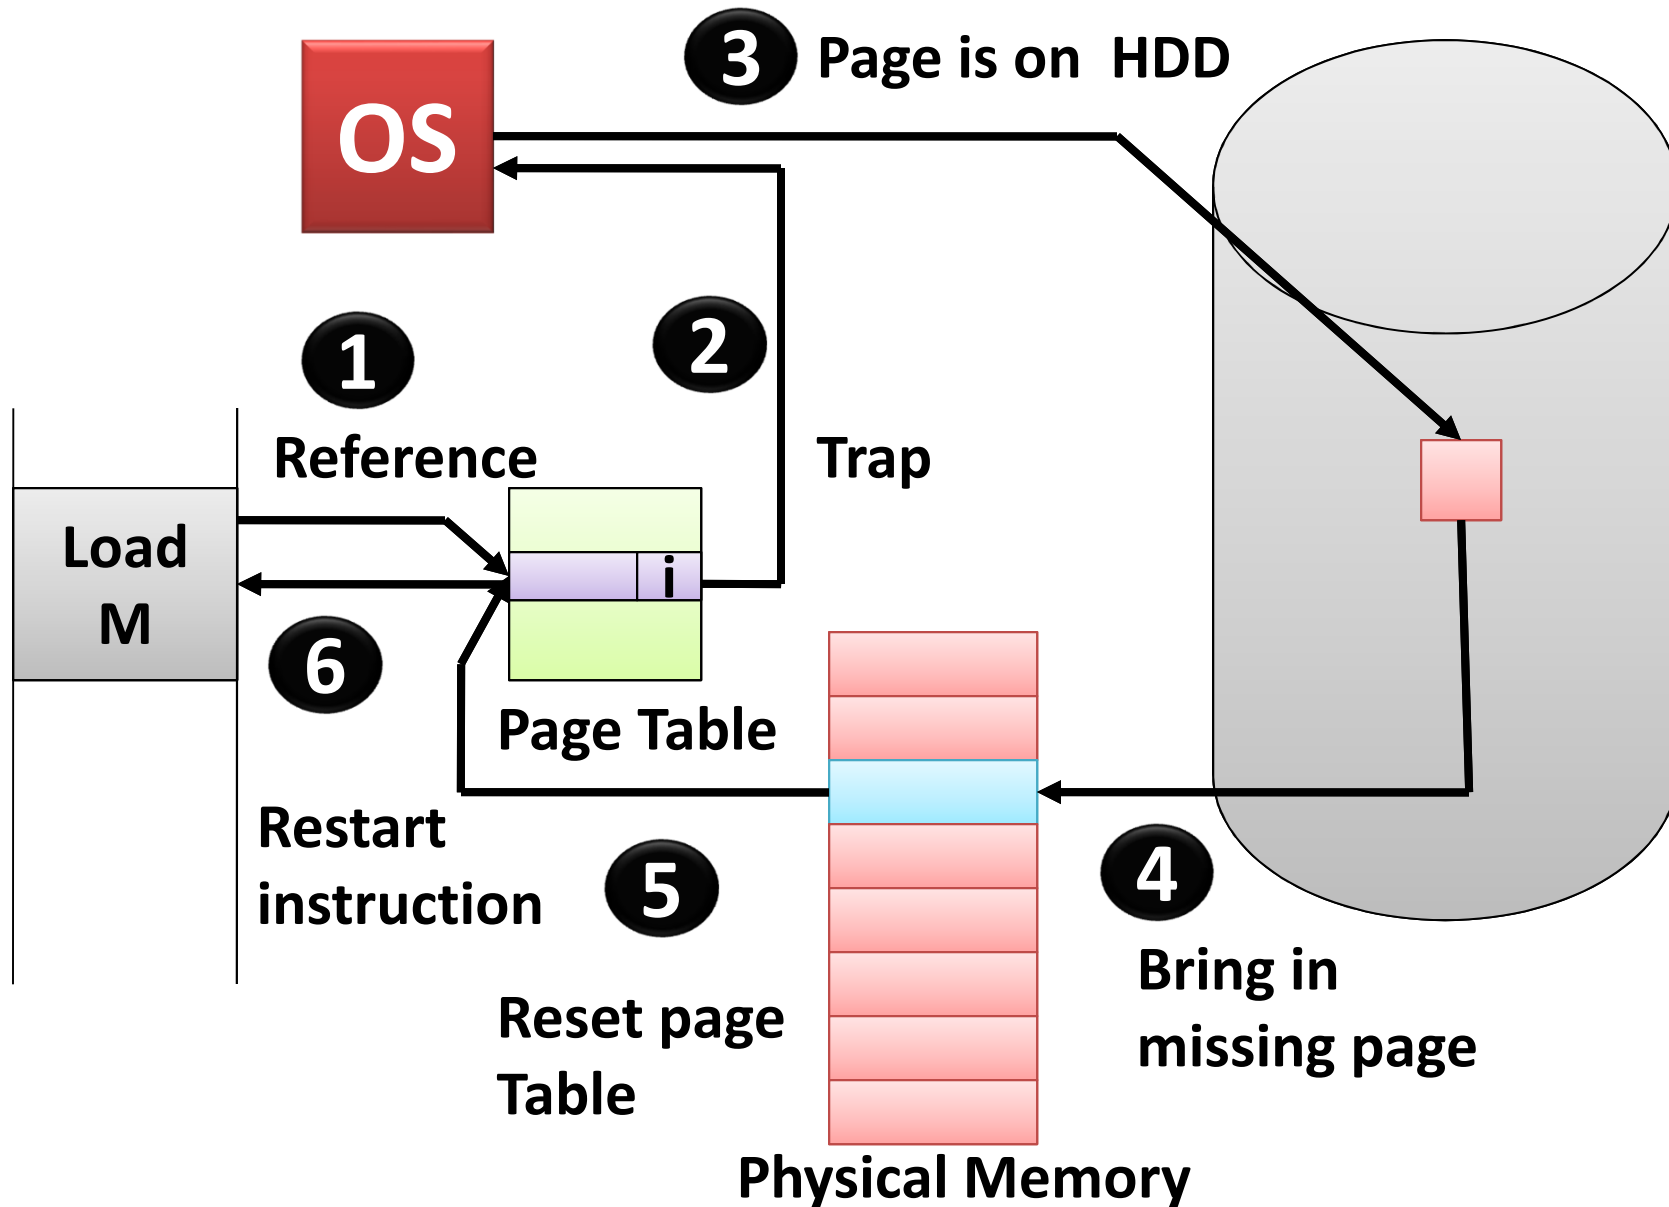


# Page Fault

- If there is a reference to a page, first reference to that page will trap to operating system:  
**page fault**
  1. Page reference
  2. OS looks at page table to decide:
    - **Invalid reference  $\Rightarrow$  abort**
    - Just not in memory
  3. Find free frame
  4. Swap page into frame via disk operation
  5. Indicate page now in memory  
Set validation bit = **v**
  6. Restart the instruction that caused page fault



# Steps in Handling a Page Fault



# Performance of Demand Paging

- Three major activities
  - Service the interrupt – careful coding means just several hundred instructions needed
  - Read the page – lots of time
  - Restart the process – again just a small amount of time
- Page Fault Rate  $0 \leq p \leq 1$
- Effective Access Time (EAT)
$$\text{EAT} = (1 - p) \times \text{memory access}$$
$$+ p (\text{page fault overhead}$$
$$+ \text{swap page out} + \text{swap page in} )$$

# Demand Paging Example

- Memory access time = 200 nanoseconds
- Average page-fault service time = 8 milliseconds
- $EAT = (1 - p) \times 200 + p (8 \text{ milliseconds})$   
 $= (1 - p) \times 200 + p \times 8,000,000$   
 $= 200 + p \times 7,999,800$
- If one out of 1,000 causes a page fault, then  
 $EAT = 8.2 \text{ microseconds.}$

This is a slowdown by a factor of **40!!**

# Demand Paging Example

- Virtual memory similar to caching in processor level
  - Except higher page fault rate is not acceptable
  - Page fault rate  $\ll$  Cache miss rate
- If want performance degradation  $< 10$  percent
  - $220 > 200 + 7,999,800 \times p$   
 $20 > 7,999,800 \times p$
  - $p < .0000025$
  - $< \text{one page fault in every } 400,000 \text{ memory accesses}$

# Demand Paging Optimizations

- Swap space I/O faster than file system I/O even if on the same device
  - Swap allocated in larger chunks, less management needed than file system
- Copy entire process image to swap space at process load time
  - Then page in and out of swap space
  - Used in older BSD Unix

# Demand Paging Optimizations

- Demand page in from program binary on disk, but discard rather than paging out
  - **Why to swap-out read only data?**
  - Used in Solaris and current BSD
  - Still need to write to swap space
    - Pages not associated with a file (like stack and heap) – **anonymous memory**
    - Pages modified in memory but not yet written back to the file system

# Demand Paging Optimizations

- Mobile systems
  - Typically don't support swapping
  - Instead, demand page from file system and reclaim read-only pages (such as code)

# Copy-on-Write

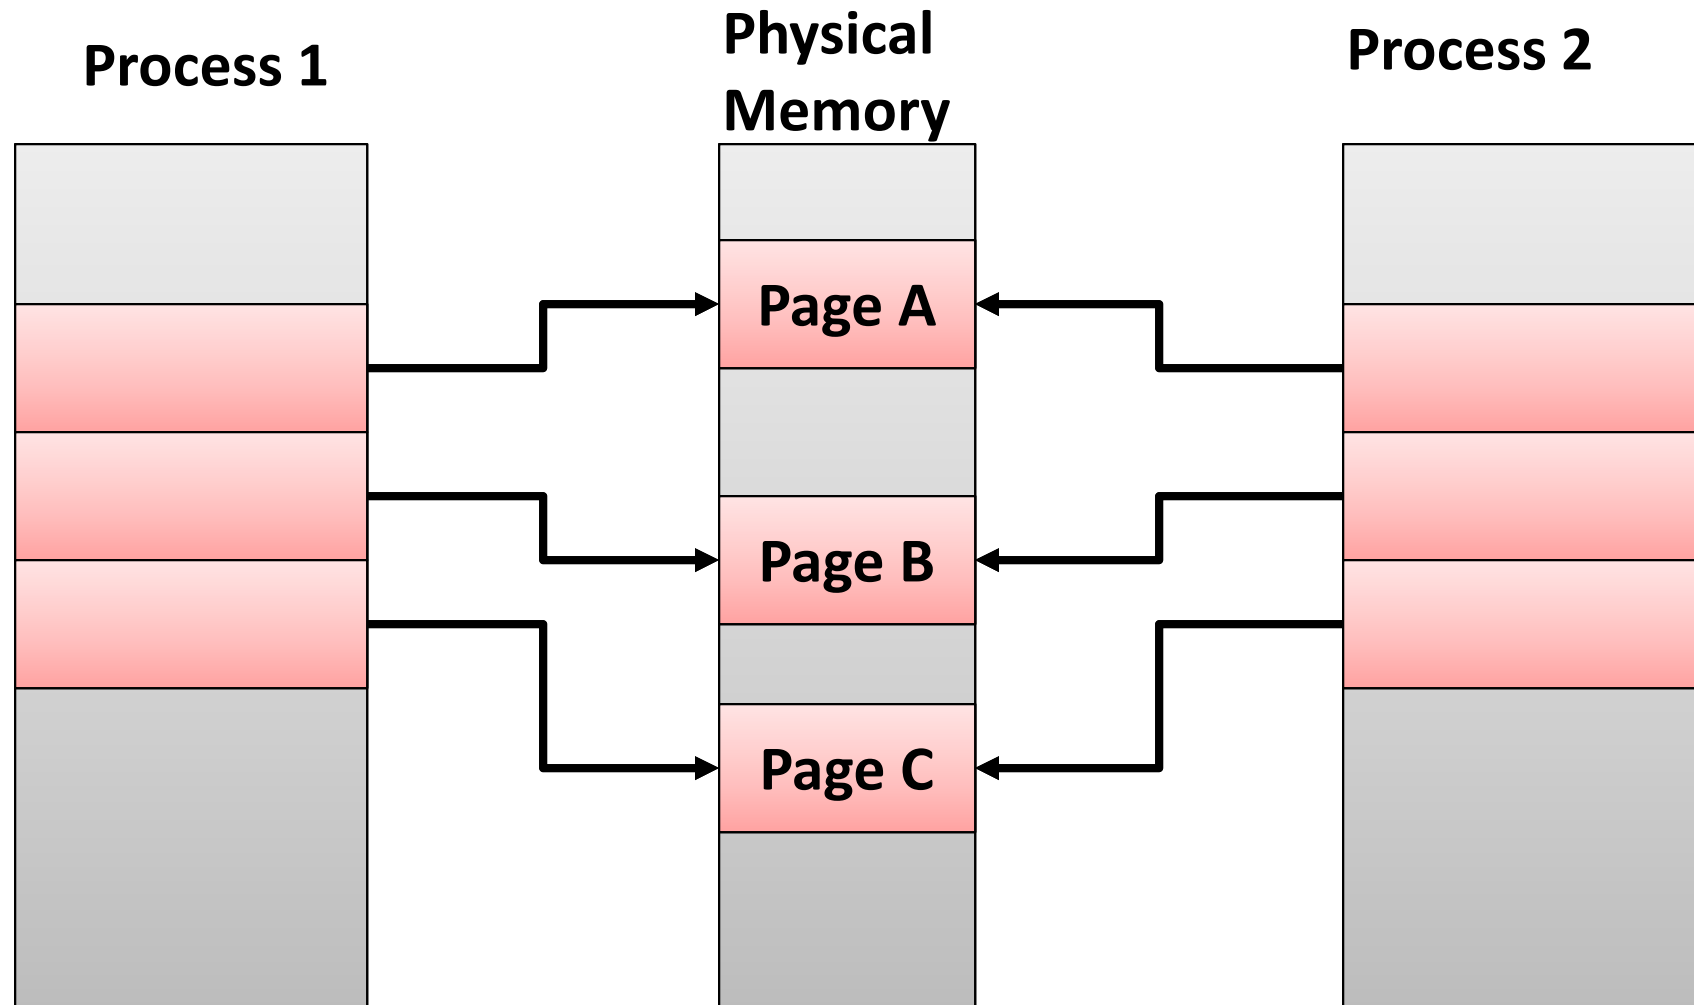
- **Copy-on-Write** (COW) allows both parent and child processes to initially *share* the same pages in memory
  - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied



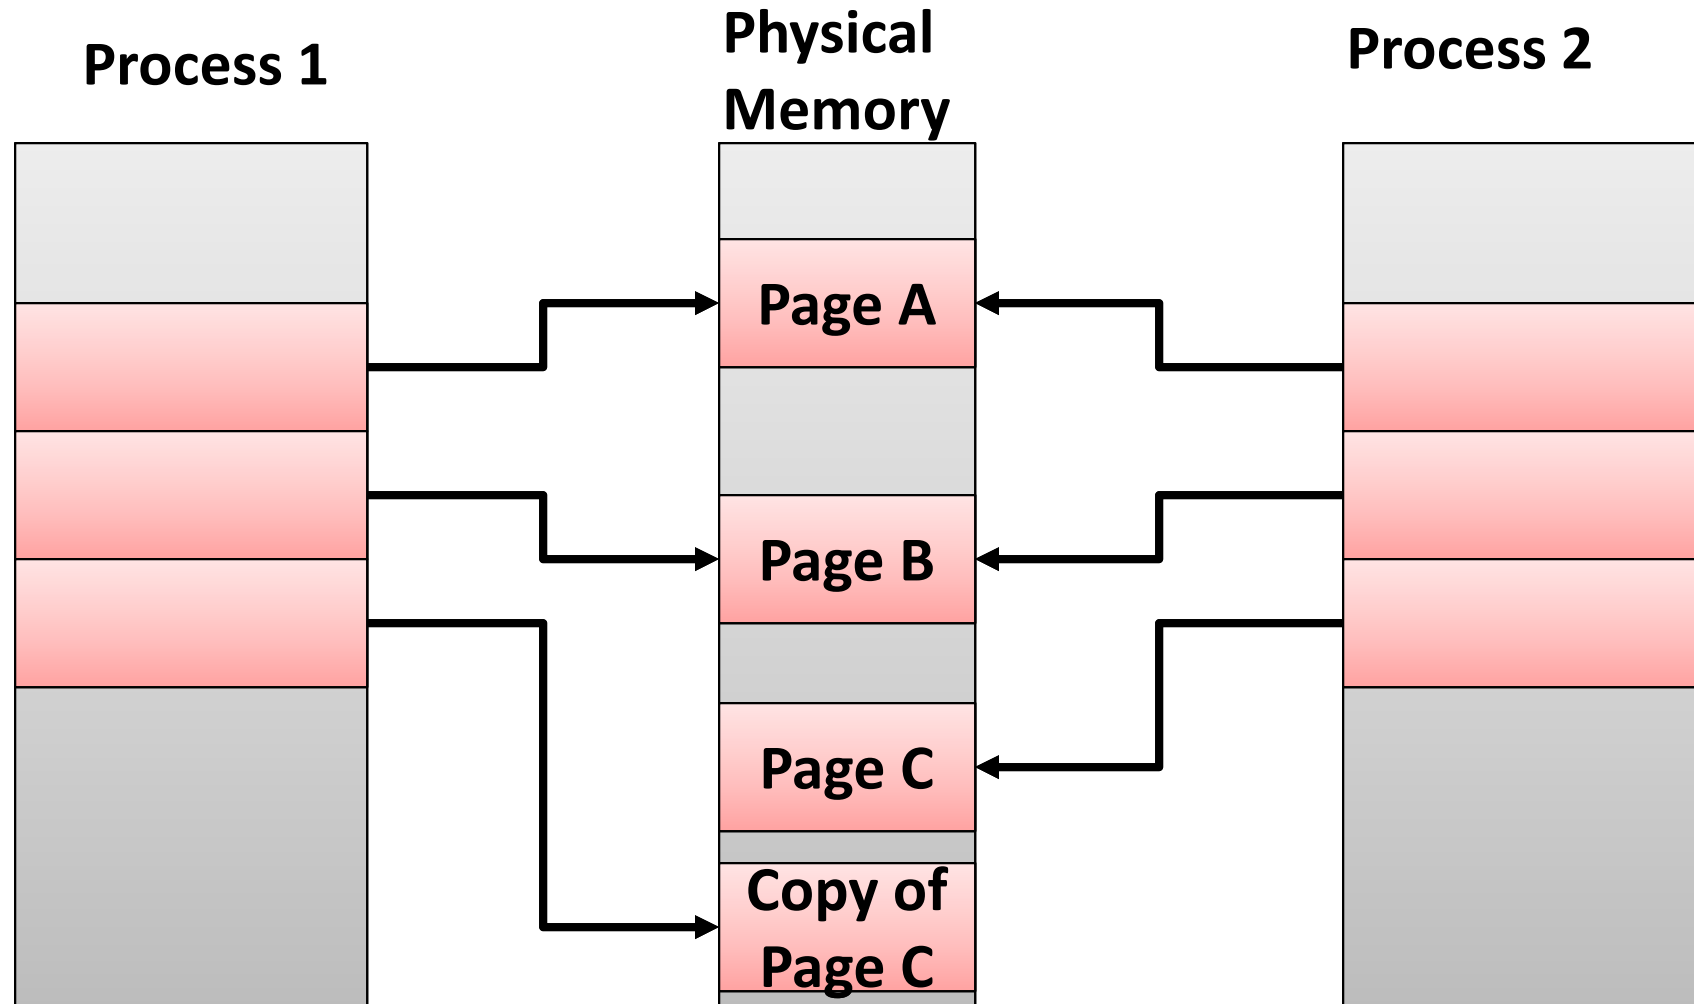
# Copy-on-Write

- In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages
  - Pool should always have free frames for fast demand page execution
    - Don't want to have to free a frame as well as other processing on page fault
  - Why zero-out a page before allocating it?
- `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
  - Designed to have child call `exec()`
  - Very efficient

# Before Process 1 Modifies Page C



# After Process 1 Modifies Page C



# What Happens if : no Free Frame?

- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- How much to allocate to each?
- **Page replacement** – find some page in memory, but not really in use, page it out
  - Algorithm – terminate? swap out? replace the page?
  - Performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

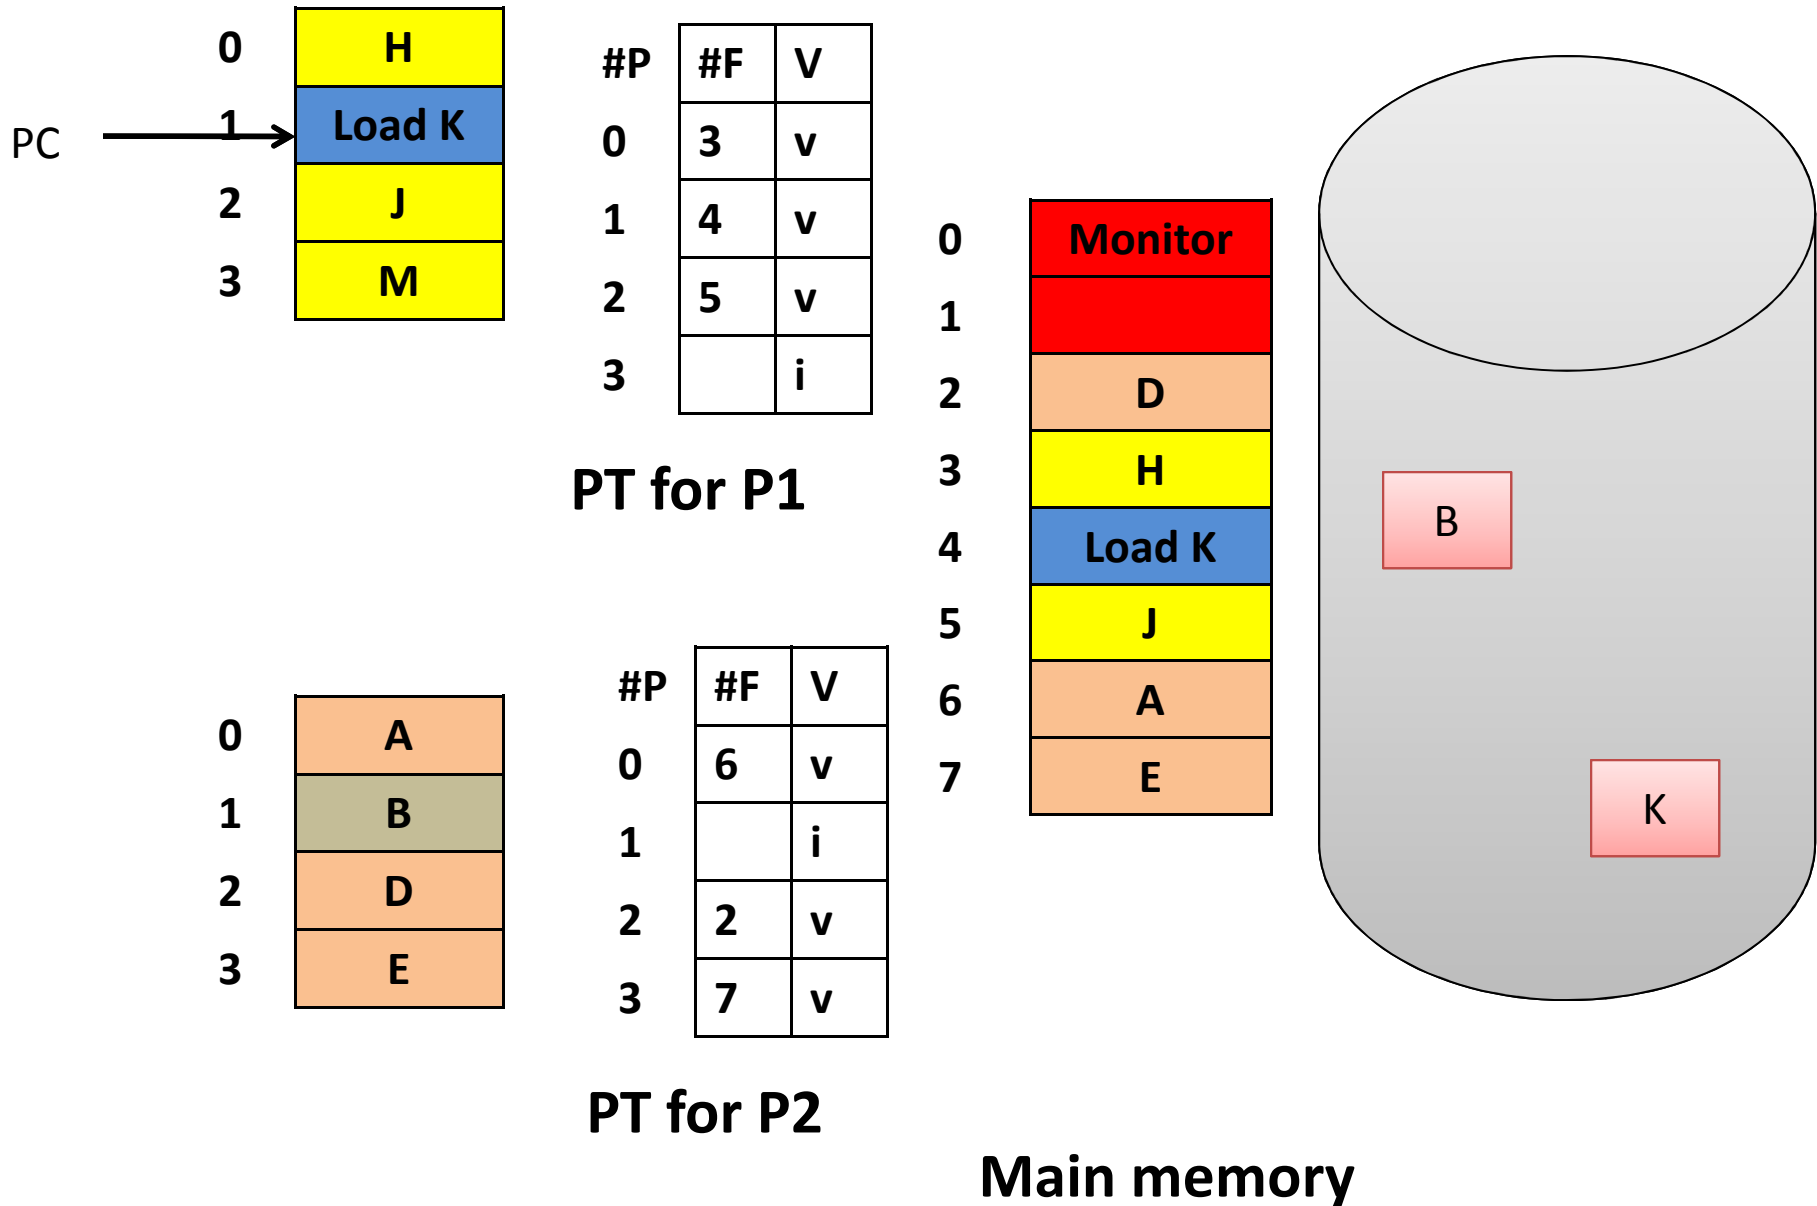
# Page Replacement

- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit**
  - Reduce overhead of page transfers
  - Only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory
  - large virtual memory can be provided on a smaller physical memory

## **Page Replacement Example**

- 2 Process : 4 logical pages per process
- 8 frame, but 2 allocated for OS
- 6 are available for user
  - Need to run 2 process of total 8 logical pages

# Need For Page Replacement

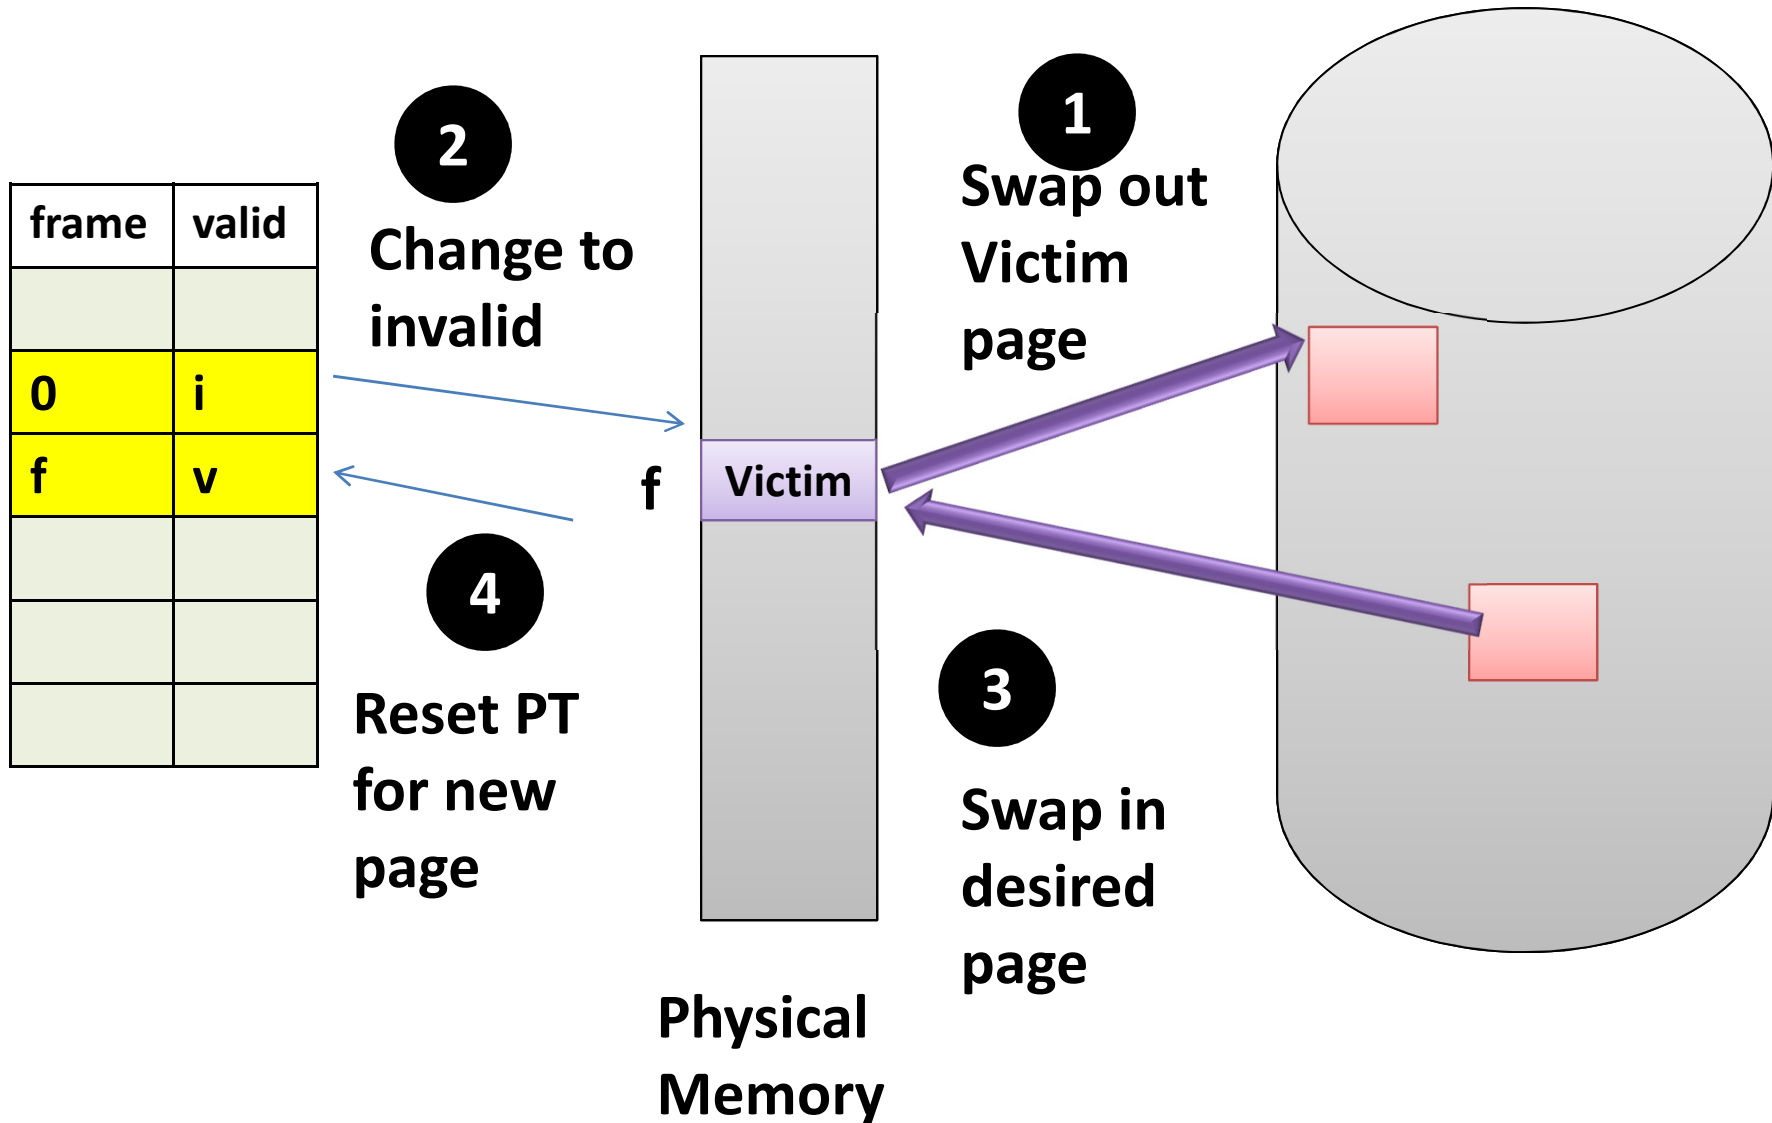


# Basic Page Replacement

- Find the location of the desired page on disk
- Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**
  - Write victim frame to disk if dirty
- Bring the desired page into the (newly) free frame; update the page and frame tables
- Continue the process by restarting the instruction that caused the trap
- Note now potentially 2 page transfers for page fault – increasing EAT



# Page Replacement



# Page and Frame Replacement Algorithms


- **Frame-allocation algorithm** determines
  - How many frames to give each process
  - Which frames to replace
- **Page-replacement algorithm**
  - Want lowest page-fault rate on both first access and re-access

# Page and Frame Replacement Algorithms

- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
  - String is just page numbers, not full addresses
  - Repeated access to the same page does not cause a page fault
  - Results depend on number of frames available
- **Example of reference string** of referenced page numbers is

**7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**

# Graph of Page Faults Versus The Number of Frames



# First-In-First-Out (FIFO) Algorithm

- Reference string:

**7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1**

- 3 frames (3 pages can be in memory at a time per process)

7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
	0	0	0	0	3	3	3	2	2	2	2	2	1	1	1	1	1	0	0
		1	1	1	1	0	0	0	3	3	3	3	3	2	2	2	2	2	1

15 page faults

# First-In-First-Out (FIFO) Algorithm

- Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
  - Adding more frames can cause more page faults!
- **Belady's Anomaly**
- How to track ages of pages?
  - Just use a FIFO queue

# Example Belady's Anomalies : with FIFO

3 2 1 0 3 2 4 3 2 1 0 4

3	2	1	0	3	2	4	4	4	1	0	0
	3	2	1	0	3	2	2	2	4	1	1
		3	2	1	0	3	3	3	2	4	4

3 frames

9 page faults


3 2 1 0 3 2 4 3 2 1 0 4

3	2	1	0	0	0	4	3	2	1	0	4
	3	2	1	1	1	0	4	3	2	1	0
		3	2	2	2	1	0	4	3	2	1
			3	3	3	2	1	0	4	3	2

4 frame

10 page faults

# FIFO Illustrating Belady's Anomaly






# Optimal Algorithm

- Replace page that will not be used for longest period of time
  - 9 is optimal for the example
- **How do you know this?**
  - **Can't read the future**
- Used for measuring how well your algorithm performs
- **Suppose you are running your program 2<sup>nd</sup> time with same data**
  - **You know what will happen in the program**


# Optimal Algorithm



# Least Recently Used (LRU) Algorithm

- Use **past** knowledge rather than **future**
- Replace page that has not been used in the most amount of time
- Associate time of last use with each page
- Use past knowledge rather than future

# LRU Algorithm



- 12 faults – better than FIFO (15) but worse than OPT (9)
- Generally good algorithm and frequently used
- But how to implement?

# LRU Algorithm Cont.

- Counter implementation
  - Every page entry has a counter;
  - every time page is referenced through this entry, copy the clock into the counter
  - When a page needs to be changed, look at the counters to find smallest value
    - **Search through table needed : Bad**

## LRU Algorithm Cont.

- Stack implementation
  - Keep a stack of page numbers in a double link form:
  - Page referenced:
    - Move it to the top
    - **Requires 6 pointers to be changed (OK)**
  - But each update more expensive
  - No search for replacement
- LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly


# Use Of A Stack to Record Most Recent Page References

Reference string

4 7 0 7 1 0 1 2 1 2 7 1 2

a b


Start



Stack before  
Time a

Singly linked list : 3 ptr  
Doubly linked list : 6 Ptr

**Move To Front  
(MTF)  
Data Structure**



Stack after  
Time b

# LRU Approximation Algorithms

- LRU needs special hardware and still slow
- **Reference bit**
  - With each page associate a bit, initially = 0
  - When page is referenced bit set to 1
  - Replace any with reference bit = 0 (if one exists)
    - We do not know the order, however



# LRU Approximation Algorithms

- Suppose every page associated with 8 bits reference bit (with a shift reg)
- When a reference made bit set to 1
- Every page reference bit shifted
- For last 8 references
  - Every page will get different value of shift reg
  - 00000000 in a Page : signifies no reference
  - 11111111 in a Page : signifies reference all the time
  - 1001010**1** is LRUed then 0010100**0**
- **All page Shift reg need to shift : Bad Part**

# LRU Approximation Algorithms

- **Second-chance algorithm**
  - Generally FIFO, plus hardware-provided reference bit and
  - **Clock** replacement
  - If page to be replaced has
    - Reference bit = 0 -> replace it
    - reference bit = 1 then:
      - Set reference bit 0, leave page in memory
      - Replace next page

# LRU Approximation Algorithms

- Second-chance algorithm with modify bit
- If page to be replaced has
  - Reference bit = 0 -> replace it
  - reference bit = 1 then:
    - Set reference bit 0, leave page in memory
    - Replace next page
- May be combined with modified bit
  - Ref bit (0) + Modified (0) ==> best guy to replace
  - 01 → not recently used but modified → not good
  - 10 => recently used but clean => probably will be used again
  - 11 → used and modified → bad candidate

# Counting Algorithms

- Keep a counter of the number of references that have been made to each page
  - Not common
- **Least Frequently Used (LFU) Algorithm**
  - Replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**
  - Based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Page-Buffering Algorithms

- Keep a pool of free frames, always
  - Then frame available when needed, not found at fault time
  - Read page into free frame and select victim to evict and add to free pool
  - When convenient, evict victim

# Page-Buffering Algorithms

- Possibly, keep list of modified pages
  - When backing store idle, write pages there and set to non-dirty
- Possibly, keep free frame contents intact and note what is in them
  - If referenced again before reused, no need to load contents again from disk
  - Generally useful to reduce penalty if wrong victim frame selected
  - **Same as Victim Buffer**

# Applications and Page Replacement

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e. databases

# Applications and Page Replacement

- Memory intensive applications can cause double buffering
  - OS keeps copy of page in memory as I/O buffer
  - Application keeps page in memory for its own work
- OS can given direct access to the disk, getting out of the way of the applications
  - **Raw disk** mode
- Bypasses buffering, locking, etc



# Page Replacement Algorithm

- Input : Reference string, number of frame
- Output: number of page fault
- **Timing information was missing in Reference string**
- **How to allocate the frame for process?**
  - Static (fixed size for process life time)
  - Dynamic

# Thanks